Flock: Forecasting How Busy a Venue Will Be, and Counting How Many People Are There

Jayden Bansal
Flock Social LLC, Pennsylvania, USA
social@flockcorp.com

October 2026. A snapshot before the mid-October retrain; to be updated with its held-out results.

Abstract

The venue card in Flock, an app for planning nights out, answers one question: how busy will this place be when the group arrives? I asked whether a venue's own recent live readings beat the popular-times chart for one venue at one hour. A serving path corrects each venue's weekly curve with a trailing offset and an age-weighted nowcast of its latest reading, with every weight fitted on earlier days. On 119,939 live readings its weights never saw (9 September to 8 October 2026), it named the exact crowd level (of five) 54.9% of the time, against 30.2% for popular times, with a mean absolute error 31.8% lower. Most of that gain is persistence: where a reading from the hour before exists, the served number is that reading, and a rule that repeats it, or else serves the corrected curve, matches the served number (within one level, 79.0% against 78.4%). The weekly curve is a poorly calibrated forecast, with observed busyness rising about half a point per point of curve, and a reading's value is spent within about an hour. The labels are a commercial provider's relative busyness readings, scaled 0 to 100 within each venue; they do not count people. To supply counts, I built Owl, a 718,550-parameter network that counts people on a 160 × 120 thermal camera in 14 to 16 ms on one Raspberry Pi 5 core. Trained on physics-based generated frames and licensed real data, it counts 82 to 96% of frames exactly on five of six real benchmark sets, where the heat-threshold rule it replaced scores 9 to 77%, and 39% on dense classroom crowds; a rule tuned to each set matches it on the two sets from Flux's own camera. Those sets also guided its development, so a sealed test and a measurement in a working venue come next. I will update this snapshot with the held-out results of a pre-registered retrain in mid-October 2026, the project's first test with every rule fixed before its data existed.

Keywords: crowd prediction; popular times; nowcasting; gradient boosting; thermal imaging; people counting; on-device inference; privacy by design

Outline. Sections 2 to 8 cover the forecast: related work, preliminary findings, methods, what goes into the model, results, discussion with the October retrain, and conclusions. Sections 9 to 18 cover Owl, the people counter. Section 19 records how both were built, Section 20 gathers the limitations, and Section 21 sets out what comes next. Appendix A lists the choices made after seeing data.

1 Introduction

1.1 Background and motivation

A group deciding where to go tonight wants one fact about each option: whether it will be busy when they arrive. The common answer in consumer maps is the venue's typical pattern for that hour of that weekday, averaged over a trailing window and drawn as a bar chart, sometimes with a coarse live indicator beside it. Möhring et al. studied the Google form of it as a behavioural signal [11]. It is a strong prior, and by construction it cannot see the specific evening: the rain, the game two blocks away, the holiday, the week the place became popular.

I build Flock, an app for planning nights out with friends. Its venue card shows one integer from 0 to 100 and one of five words: *Quiet*, *Not Busy*, *Steady*, *Busy* and *Packed*. Whether that number is right decides whether a plan works. This project measures how much better than popular times that number can be, and what it takes.

1.2 Research question

On the venue-hours a production system actually serves, can a system that watches venues over time beat each venue's own typical pattern, and by how much?

1.3 Hypotheses

- **H1.** Correcting the popular-times curve with the venue's own recent live readings reduces error against held-out live readings, on both the point scale and the five-level scale.
- **H2.** A model that predicts how an hour departs from the curve, using context such as weather, events and calendars but no recent history, beats the curve on cities it has never seen.
- **H3.** The value of a live reading decays within hours, so the weight it deserves falls with its age.

1.4 Contributions

The main finding is that the freshest live reading is worth more than any correction to the venue's curve that this project found. Over the 30-day test, adding the latest reading lifts within-10 accuracy from 31.2% to 55.2% (Section 6.5). Where the venue was read an hour earlier, the served number is that reading, and it lands on the exact crowd level 77.0% of the time (Section 6.3). A rule that uses only an hour-old reading, and otherwise the corrected curve, matches the served number (Section 6.1). A venue's departure from its curve correlates 0.74 from one hour to the next over the test month and 0.26 across two, so a reading's value is spent within about an hour (Section 3.2). Most of the gain over popular times is persistence, and persistence of the latest reading is the baseline a curve correction has to beat. The contributions are:

- A serving path, in production since 26 September 2026, that adds a trailing per-venue offset and an age-weighted nowcast to the venue's curve, with every weight fitted strictly before the rows that score it. On 119,939 live readings its weights never saw, it names the exact crowd level 54.9% of the time, against 30.8% for the configuration it replaced and 30.2% for popular times, and cuts the

popular-times baseline's MAE by 31.8% (Section 6.1). Scored beside persistence baselines, most of that gain is the latest reading itself.

- Three properties of typical-week curves, measured on this provider's readings: as a forecast the curve is poorly calibrated, with observed busyness rising about half a point per point of curve; a venue's departure from its curve stays correlated for one to two hours; and within-10 accuracy partly rewards predictions that land on the labels' 5-point lattice (Section 3).
- A gradient-boosted delta model, evaluated on three cities it never trained on and on the served population, that takes R^2 above zero at single-venue, single-hour granularity (Section 6.2), and a complete account of what goes into it (Section 5).
- An evaluation apparatus that outlasts any one model: a ship gate against production as served, trainer-server parity tests, a dataset that records its own provenance, and a retrain test whose rules were committed before its data existed (Sections 4 and 7.2).

1.5 The second study: counting the room

The forecast target in Sections 3 to 8 is a busyness score normalised per venue: 100 means the busiest that venue gets, and the score does not count people. Sections 9 to 18 describe Owl, an on-device people counter for Flux, Flock's venue sensor, which counts the people a thermal camera can see and sends a few numbers, never an image. Section 19 records how both systems were built, Section 20 gathers the limitations of both studies, and Section 21 sets out the next measurements: the October retrain, and joining the two studies so that the forecast is checked against real counts.

2 Related Work

Point-of-interest visit forecasting. The closest published task forecasts hourly visits to named points of interest. BysGNN [7] forecasts hourly visits to 400 POIs in each of five US cities from SafeGraph data, whose hourly visits also underlie mobility network models of epidemic spread [3]. It reads the previous 24 hours of visits at every venue in a city and forecasts the next six hours, a live feed a consumer app rarely holds. Against a Historical Average baseline (the same weekday over the previous month) that its authors liken to the popular-times graph, its MAE is 40% to 54% lower across the five cities at its 400-venue scale (Houston 8.860 to 4.095); against last week's value at the same hour the gain is 14% to 19% (both computed from its tables). Its more often quoted 4.34% to 6.71% margin is over the strongest competing graph model. Its MAPE of 60% to 80% shows what precision the task admits; MAPE itself degenerates near zero [8]. Timokhin et al. [16] estimate hourly venue popularity from venue attributes and compare popular-times data with WiFi counts.

Occupancy prediction and granularity. Bollenbach et al. [2] report R^2 up to 0.87 for whole-beach visitor counts one step ahead, while section-level occupancy forecasts one to three days ahead have R^2 near zero and often below it (−1.26 to 0.42 across the best models). Following Moore et al. [12], they read R^2 above 0.3 as weak, 0.5 as moderate and 0.7 as strong. My task sits at the granularity a crowd card answers: one venue and one hour.

Beating the trivial baseline is not guaranteed. In a beach-occupancy study [1], a feature-based random forest (MAE 94) and SARIMA (MAE 82) both had higher MAE than naive persistence (MAE 79), though both had lower RMSE. In weekly foot-traffic forecasting, a simple weighted average beat an LSTM [9]. These results are why every candidate in this project is scored against the number production serves and against persistence of the latest reading, the standard naive benchmark [8], on the same rows (Sections 4.4 and 6.1).

3 Preliminary Findings

Before designing anything, I measured the popular-times baseline against provenance-verified live readings. Each finding below fixed one design decision.

3.1 The curve is a poorly calibrated forecast

I joined every verified live reading to its own venue's curve value for the same weekday and hour, giving 9,903 paired rows on which the curve can be compared with the provider's live reading for that hour. The gap inverts with the level of the curve (Figure 1). In hours whose curve value is at most 25% of the venue's daily peak, observed busyness sits 15.60 points *above* the curve (1,162 rows). In hours above 75% of the peak it sits 13.56 points *below* (5,746 rows). The middle quarters deviate by +1.68 and −4.99.

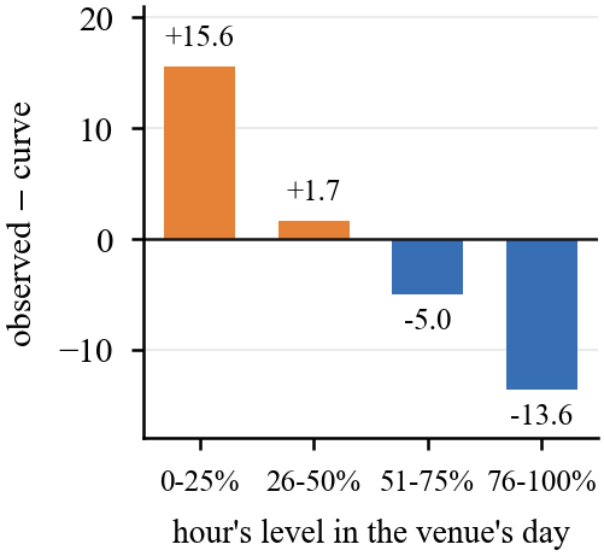


Figure 1. Observed minus curve, by the hour's level in the venue's own day (9,839 of the 9,903 paired rows).

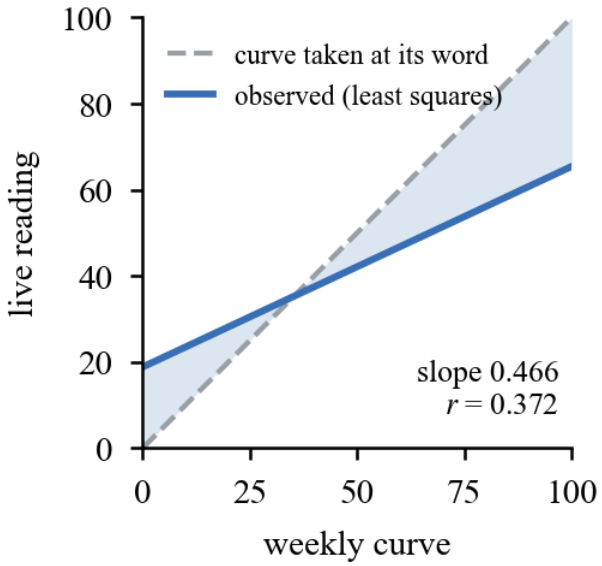


Figure 2. The least-squares fit of live readings on the curve (9,903 rows), against the identity line.

Regressing observed busyness on the curve gives *observed* = 18.82 + 0.466 × *curve*, with $r = 0.372$ (Figure 2): each point of curve is worth about half a point of observed busyness, and the curve explains about 14% of the variance in live readings. The readings themselves are more spread out than the curve (standard deviation 33.2 against 26.4), and the slope is low because the correlation is low. The best

forecast from such a curve shrinks it toward the mean, the regression effect in prediction [5, 6]. The 30-day test window repeats the pattern: *observed* = 18.62 + 0.497 × *curve*, r = 0.360, with quiet hours 11.5 points above the curve and busy hours 13.9 below. The same effect explains why venues' offsets averaged about −11 (1,048 venues, 2026-09-06), which I had first read as a plain bias; on these 9,903 rows the mean deviation is −6.3. **Design decision:** a correction must be able to move quiet hours up and busy hours down, which one global offset cannot do.

3.2 A venue's recent state persists, then fades

In the first week of hourly collection (1 to 8 September) the provider's live value was sticky: 80.9% of consecutive-hour readings of a venue were identical, and a venue's deviation from its curve correlated 0.815 across one hour, 0.487 across two and 0.203 across three (Figure 3). Used as the prediction for the next hour, a reading one hour old landed within one crowd level 92.0% of the time and within 10 points 85.9% (3,165 readings). Nothing in a venue's typical week carries this information. **Design decision:** the latest reading should enter the served number directly, at a weight that decays on the same timescale.

The stickiness did not last. Identical consecutive readings fell to 69.3% from 9 to 25 September and to 44.8% and 47.5% in the two periods after, and over the 30-day test window the deviation correlates 0.739, 0.255 and 0.158 at one, two and three hours. Whether the early stickiness came from how often the provider refreshed its value or from when the collector read it cannot be told from these data.

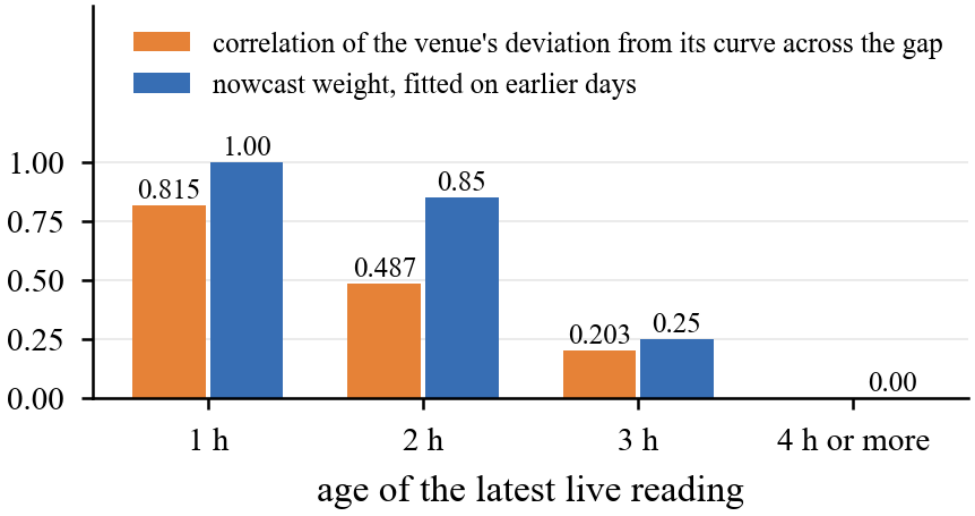


Figure 3. How a venue's deviation from its curve persisted over 1 to 8 September, against the nowcast weights fitted on 1 to 5 September (Section 4.4). The fitted weights fall with age, more slowly than the correlation at two hours. Over the 30-day test the correlation at two hours fell to 0.26.

3.3 The within-10 metric rewards the grid

The provider's labels lie on a lattice of multiples of 5, and the stored weekly curve is the provider's typical-week value as given, one per venue, weekday and hour, so it sits on the same lattice. A forward split shows what that does to the metric. I fitted each venue's offset on its earlier readings and scored its later ones (894 venues, 3,930 later readings) at offset weights λ from 0 to 1 (Figure 4).

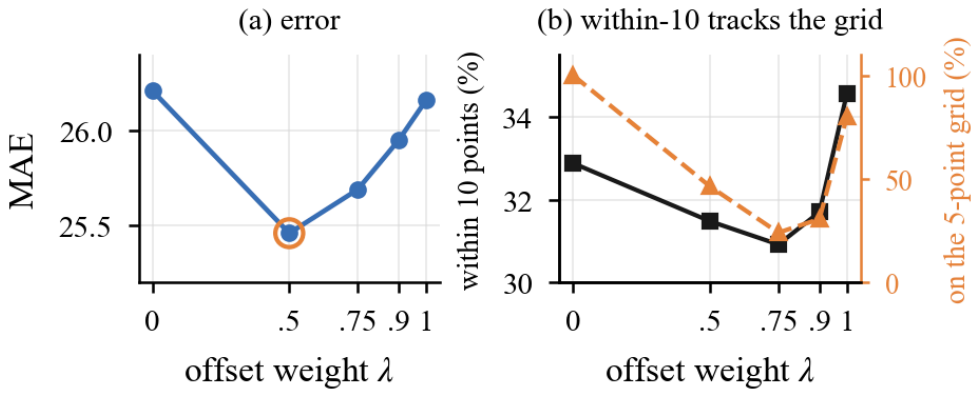


Figure 4. The forward offset test. (a) MAE is best at $\lambda = 0.5$. (b) The within-10 rate follows the share of predictions on the 5-point grid, not the error.

MAE is best at $\lambda = 0.5$ (25.46, against 26.21 with no offset). Within-10 is highest at the two weights whose predictions most often land on the grid: $\lambda = 1$ (34.55%, 79.9% on-grid) and $\lambda = 0$ (32.88%, 100% on-grid). The median error is 20 at every weight. On a lattice-valued dataset, within-10 partly measures grid coincidence. **Design decision:** where a gridded and a non-gridded predictor are compared on within-10, the comparison reports the on-grid rate or projects both to the lattice (Section 6.4); production's within-10 margins in Table 5, 24 points over the weekly curve and over curve plus offset, are far larger than the within-10 variation across the offset weights tested here.

4 Methods

4.1 Definitions and problem statement

Busyness and crowd levels. For a venue v and an hour h, busyness $y \in [0, 100]$ is a relative score normalised per venue, so that 100 is the busiest that venue gets; it is not a count of people. The app publishes round($\hat{y}$) and a word: *Quiet* (0 to 20), *Not Busy* (21 to 39), *Steady* (40 to 69), *Busy* (70 to 84) and *Packed* (85 to 100).

The weekly curve. Let $b(v, d, h)$ be the venue's typical value for weekday d and hour h: the venue's stored weekly value on its local clock, smoothed with its two neighbouring hours at weights 0.6, 0.2 and 0.2. The served number starts from this curve. The popular-times baseline is the stored weekly value itself, before smoothing, as a popular-times chart shows it. The smoothed curve is a whole number but not always a multiple of 5.

Problem. Given a venue v, a target hour h, the curve b, and the live readings of v taken strictly before h, produce $\hat{y}(v, h) \in \{0, \ldots, 100\}$ and its crowd level, to minimise error against the live reading at h. The target hour's own reading is never an input.

4.2 Data collection

Table 1. Data sources. The busyness provider is a commercial venue-busyness service; its scores are relative, 0 to 100 within each venue.

Source	Contribution
Weekly curves	The busyness provider's typical week for each venue: 168 hourly values on a 21-point lattice of multiples of 5. These rows define the curve.
Live readings	The busyness provider's reading for one venue at one hour on a real date, collected with a provenance label (live or provider forecast) since 2026-09-01, hourly from 3 September. About 1,380 Pennsylvania venues are scheduled every hour and Miami venues at two fixed hours a day. The 30-day window holds about 7,600 model-served live readings a day from 26 September to 2 October and about 4,600 a day from 3 to 8 October, after a drop on the provider's side.
Weather	Temperature, humidity, wind, condition code and rain flag at the hour, from OpenWeather.
Events	Nearest ticketed event's distance, type and attendance, and counts within 2 km, from Ticketmaster listings. For the October retrain, 627 professional and college games from TheSportsDB, refreshed daily.
Venue attributes	Category, rating, review count, price level and up to three place types, from a places listing. The venue's location is used to look up weather, daylight and neighbours and is never a feature.
Calendars	Holiday, school break, holiday eve and special nights across 29 country and city scopes.
Owner readings	A venue's own live 0 to 100 report, at training weight 0.10.
In-app reports	Users' crowd reports, verified and aggregated per venue over 28 days (the feedback family of Table 3); none existed when v2.6.0 was trained.

Procedure. Table 1 lists every source. An hourly collector reads every tracked venue's live value and stores one row per venue per hour, stamped with the venue's local clock and a provenance label. Since 2026-09-25 the full sweep finishes inside the hour (8 min 42 s for a 5 a.m. sweep that called 358 venues and skipped the rest as closed, about 31 minutes for a 1,321-call evening sweep); before that it stopped at its 50-minute budget. Weekly curves and context tables are refreshed on their own schedules.

Size. At the training run the corpus held 3,912,357 rows: 3,454,955 weekly and 457,402 realtime, across 34,784 venues in 34 cities, collected between 2026-03-10 and 2026-05-18. Realtime rows were scarce: 457,402 observations over 18,442 venues is 25 per venue across 168 weekly slots (26.3 on 2026-09-06: 466,179 rows over 17,722 venues after the 2026-09-05 duplicate-venue merge).

Integrity. Unique keys enforce one weekly row per (venue, weekday, hour) and one realtime row per (venue, date, hour). Before they existed, I found and repaired three faults: the hour column was an array index in 3,536,360 rows (slot 18 held the venue's midnight), there were 491,100 duplicate rows, and 63% of rows carried an impossible month-season pair. Every accuracy figure in this paper comes from after the repair. The 457,402 realtime rows collected before provenance was recorded cannot be separated into live readings and provider forecasts, and stay marked *unknown*. "Live" in this paper means the provenance-verified readings collected since 2026-09-01 (hourly from 3 September).

4.3 Variables and controls

Table 2 lists the variables and the controls that hold every comparison fixed.

Table 2. Variables and controls.

Type	In this project
Independent variable	The serving configuration or model that produces the number: weekly curve, curve + offset, curve + offset + nowcast, two persistence baselines, rule engine, delta model v2.6.0, incumbent v2.5.0.
Dependent variables	Within one level, exact level, band MAE (the mean number of levels between predicted and observed), MAE, R^2, within 10 and within 15 points, bias.
Controlled: the rows	Every configuration is scored on the identical held-out readings, rebuilt through the production serving code.
Controlled: time	The offset and nowcast weights are fitted on 2026-09-01 to 09-05. The main evaluation covers 2026-09-09 to 10-08; the first test covered 09-06 to 09-08.
Controlled: geography	Miami, Tokyo and Barcelona never enter the shipped delta model's training, and each learner-comparison fold leaves out the city it scores. Temporal tests use a venue's earlier readings, never the scored reading or a later one.
Controlled: population	Scored rows must satisfy the serving predicate (realtime, curve above zero), one function shared by trainer and evaluator.
Reference	A constant "Not Busy", which shows how much of any level-agreement rate a do-nothing answer earns.

4.4 System design

Figure 5 shows the system. Three data sources feed two pipelines. The online pipeline serves every request. The offline pipeline trains a delta model, which can replace the online path only by passing a two-part ship gate.

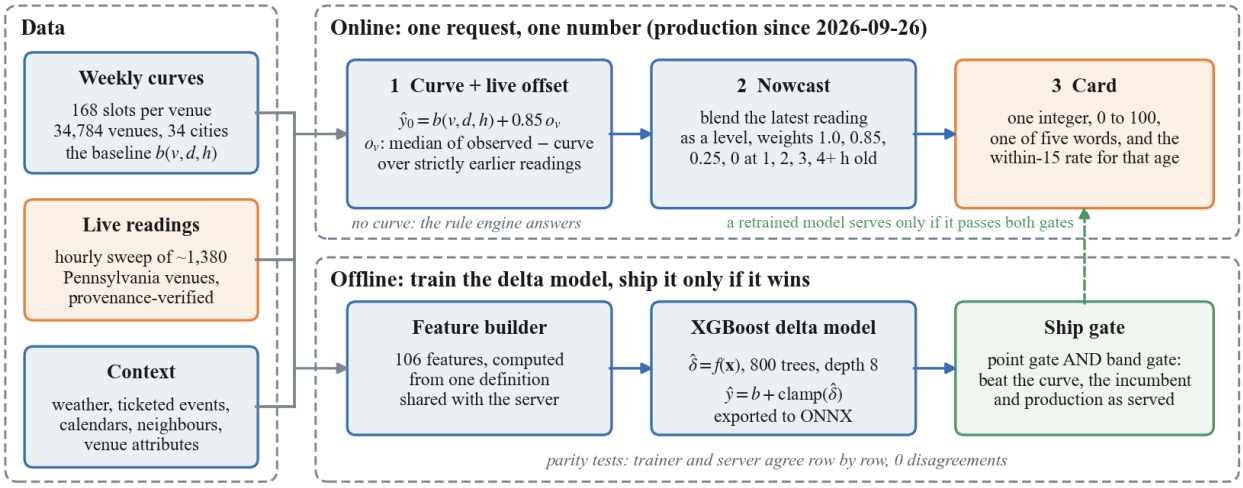


Figure 5. The system. One data column feeds an online serving path (in production since 2026-09-26) and an offline delta model, which serves only if it passes both halves of the ship gate.

Online correction. For each venue, let o_v be the median of (*observed* − *curve*) over the venue's newest 20 live readings taken strictly before the hour being scored and within 28 days, clamped to ±50 and used only with at least two readings (otherwise 0). Let r_v be its most recent reading from an earlier hour, and a the age of that reading in hours. The served number is

$$\hat{y}_0(v,\, h) = b(v,\, d,\, h) + 0.85 \cdot o_v, \tag{1}$$

$$\hat{y}(v,\, h) = (1 - w_a) \cdot \hat{y}_0(v,\, h) + w_a \cdot r_v, \tag{2}$$

with $w_1 = 1.0$, $w_2 = 0.85$, $w_3 = 0.25$ and $w_a = 0$ for $a \geq 4$. The offset weight and all four blend weights were fitted by MAE on 2026-09-01 to 09-05. At an age of one hour the fitted weight is 1.0, so production then serves the last reading unchanged: the "repeat the last reading" baseline of Section 3.2 is the case the fit chose for fresh readings, and the offset takes over as readings age. The offset weight of 0.85 is the value this fit chose on those five days; the forward test of Section 3.3, on a different split, preferred 0.5. The weights were fitted while readings were at their stickiest, and Section 6.3 shows that over the month the two- and three-hour weights no longer helped. The card also publishes a confidence: the within-15 rate measured for the number's own reading age (Section 6.3). Section 5 lists every step of the serving path and every input it reads.

Offline delta model. The model predicts how an hour departs from the curve, $\hat{\delta} = f(\mathbf{x})$, from 106 features in eleven families, listed in full in Table 3: time and calendar, daylight, holidays and school breaks, weather, ticketed events, category curves, neighbouring venues' curves, venue attributes, place types, in-app feedback and a row-type flag. The venue's own curve is the anchor and is refused as an input. Every column that would leak the label is excluded, following Kapoor and Narayanan's taxonomy [10]. The learner is XGBoost [4]: 800 trees at depth 8, learning rate 0.01, subsample 0.9, column subsample 0.8, minimum child weight 7, L1 0.5, L2 1.0. The score is rebuilt by a fixed chain, clipped to [0, 100] and rounded:

$$\hat{y} = \text{round}(\text{clip}(b + 0.5\ o_v + \text{clamp}(\hat{\delta}, \pm 50), 0, 100)). \tag{3}$$

On the geographic holdout no venue has earlier readings, so $o_v = 0$ there. The v2.6.0 ship gate scored the delta with a ±30 clamp and without rounding (Table 6); Section 6.2 also gives the model under Equation (3) exactly.

Ship gate. A candidate model must pass a point gate and a band gate.

- *Point gate:* on the held-out cities, beat the curve by 5 MAE or 0.10 R^2, not regress MAE, and match the incumbent's own within-10 on the same rows.
- *Band gate:* on at least 1,000 held-out live readings over 5 or more dates, it must:
  - beat production *as served* on within one level (the date-block 95% lower bound above −1 point);
  - not be worse on within-10 (the date-block 95% lower bound above −1 point);
  - beat the weekly curve on within one level (lower bound above 0) with band MAE no worse;
  - not fall below the rule engine on within one level;
  - hold every city with 300 or more readings within 2 points;
  - raise MAE by at most 1 point.

The server fails closed: an artifact that did not pass cannot be loaded.

4.5 Evaluation protocol and statistics

- **Time holdout.** The served number is scored on every model-served live reading from 2026-09-09, the day after its first test, to 2026-10-08, the day of the export: 119,939 readings over 30 dates at 1,754 venues (Philadelphia 63,225, Lehigh Valley 49,503, Miami 7,211), each rebuilt through the production serving code. Two persistence baselines are scored beside it: (a) the newest earlier reading if at most 12 hours old, else popular times; and (b) the reading from the hour before if there is one, else curve plus offset. The offset and nowcast weights were fitted on 2026-09-01 to 09-05 and have not changed since, so no reading in the window informed them. The first test, 4,183 readings from 2026-09-06 to 09-08 (Lehigh Valley 3,615, Miami 568), is reported beside it.
- **Geographic holdout.** The delta model is scored on the *gate slice*: the 67,249 realtime rows with a usable curve in the three held-out cities, a whole-city holdout [15]. Section 6.2 explains why these cities act as a validation set.
- **Cross-validation.** Leave-one-city-out over the 30 training cities, the block design recommended for spatially structured data [15]. The corpus's 34th city, Beijing, holds no rows in v2.6.0's training frame.
- **Reproduction.** Every geographic-holdout number is recomputed from the saved model and holdout by scripts that pin seeds and thread counts, and those scripts also produce the intervals and the learner comparison (Tables 8 and 11).
- **Statistics.** Live-reading intervals are date-block bootstraps that resample whole observation dates (2,000 resamples) [15]; geographic-holdout intervals use the two-way cluster bootstrap of Table 8.

Tables 4 and 5 are computed from the per-reading rows of the 30-day window, so each of their intervals resamples 30 date blocks. Readings of one venue on one night are correlated, and the date blocks carry that: on the exact-level margin over popular times they widen the interval about 4 times over treating readings as independent, about 15 times the variance. A two-way bootstrap that resamples venues and dates together widens the intervals by a further 10% to 17%; where it changes a conclusion, Table 5 says so. In the results tables, level agreement is reported with exact level and band MAE beside it, because a constant scores nearly as high on it (Section 6.1).

4.6 Reproducibility safeguards

Following common reproducibility practice [14], training pins the device, thread count and seed, and a test confirms that two runs on the same data produce bit-identical predictions. The model is exported to ONNX and served in-process through ONNX Runtime [13]. Each feature family is computed by the trainer and the server from one definition, and a test requires them to agree row by row. All report zero disagreements:

- the neighbour arithmetic, on 8,006 readings;
- the live-reading features;
- the game features;
- 62 calendar, weather and event columns;
- the event collection, on 400 listings.

An offline replay used for every live score is proven by a test to publish exactly what the serving function publishes.

5 What Goes into the Model

5.1 Two ways the number is made

Two pieces of arithmetic can produce the number on a crowd card, and they read different inputs. Since 26 September 2026 production has served the *online correction* of Equations (1) and (2): the venue's own weekly curve, its trailing offset and its latest earlier live reading. In that mode no model runs, so weather, events, calendars and neighbours do not reach the number. The *model mode* is the code's default when no switch is set, and it is the path on which the delta model v2.6.0-starling runs. That model takes exactly 106 inputs, listed by family in Table 3. Both modes share the same fallback: a venue with no curve, or a server with no loaded model, is answered by a rule engine.

The shipped model was trained on 1,934,988 rows from 30 cities, with weekly rows at weight 0.05 and spring realtime rows (provenance unknown) at 1.0, and with coordinates, venue identifiers, the venue's own curve and every provenance column excluded.

Table 3. Every input of the delta model v2.6.0-starling (106 features), by family, with its source at serving time, refresh rate and use in the shipped trees. *Splits*: the family's share of the 167,758 decision splits in the 800 trees; *roots*: trees that split on the family first; both counted from the shipped ONNX file. The last three rows are not inputs of

v2.6.0: the venue's curve is the anchor, and the families marked † are built for the October retrain.

Family	Features	Source at serving	Refresh	Use in the shipped trees
Time (19)	hour, weekday, month; sine and cosine of each; weekend; Friday or Saturday after 18:00; lunch, dinner, late night, morning; four season flags	The venue-local clock of the hour being scored	Every request	12.4% of splits; 349 roots, 274 of them on month. Split gain: month 20.1%, its sine and cosine 4.8%
Daylight (3)	daylight hours, hours after sunset, after-sunset flag	Solar estimate from the venue's latitude and the month	Every request	11.2% of splits; 22 roots. Daylight hours alone 6.6%
Holidays (6)	holiday, school break, holiday eve, special night, special-night boost and suppress weights	Ten US federal holidays a year, US school-break windows, and 31 national holiday calendars with 1,351 dated special nights over 29 scopes (2025 to 2028)	Regenerated yearly	1.1% of splits; 55 roots, all on the suppress weight
Weather (18)	temperature, humidity, wind, rain flag; nine condition groups; temperature anomaly against a latitude-band by month climate table; four weather interactions	Current conditions keyed to about 1 km; a missing reading becomes the climate norm	Cached 30 min per location	19.5% of splits; 140 roots. Humidity, anomaly, temperature and wind each 3.8% to 4.6%
Events (16)	any event within 2 km; nearest event's attendance and its log, and distance; its type as five flags (arts, family, music, sports, other); count, total attendance and its log within 2 km; large event; event × weekend, dinner, bar	Ticketed events under way at the scored hour within 2 km; a failed lookup enters as no event	Cached 1 h	1.3% of splits; 5 roots
Category curves (2)	typical value of the venue's category at this weekday and hour; the same, refined by price and rating tier	Tables fitted on the training corpus, shipped in the model's metadata	Fixed until retrain	9.5% of splits; 17 roots
Neighbours (2)	log count of venues with a curve within about 830 m; their mean curve value at this weekday and hour, the venue itself excluded	Stored weekly curves of nearby venues	Cached 24 h	13.8% of splits; 33 roots. Second and third most split-on features (7.1% and 6.7%)
Venue (5)	category (13 classes), price level, rating, review count and its log	The venue's places listing; missing values take the corpus median, or 0 reviews	Every request	20.0% of splits; 95 roots. Log review count is the most split-on feature (9.3%)
Place types (30)	one flag for each of the 30 most common place types, from the venue's first three types	Same listing	Every request	9.4% of splits; 84 roots
Feedback (4)	mean crowd report, log report count, has-report flag, mean error of past predictions	Per-venue aggregate of verified in-app reports from the last 28 days	Cached 1 h	Never split on: no verified report existed at training time

Family	Features	Source at serving	Refresh	Use in the shipped trees
Row type (1)	realtime or typical-week row	Always 1 at serving	Constant	1.9% of splits, 0 roots; split gain 6.1%; dropped for the retrain
Own curve (0)	refused as an input	The busyness provider's typical week, smoothed 0.6 / 0.2 / 0.2	Rebuilt after each hourly sweep	The model predicts the departure from this curve, and the served number adds that departure back
Live readings† (5)	last earlier reading's departure from the curve and its age; trailing offset and its count; the curve one hour earlier	The venue's own live readings strictly before the scored hour	Hourly	47% of a candidate's split gain in the September rehearsal
Sports games† (6)	a tracked team plays today; number of games; evening game; home game; distance to the home arena; arena within 10 km	Public sports schedules, for venues within 60 km of a tracked arena	Daily	0.02% of split gain in the rehearsal

Reading the last column. The record holds split gain for the month features (month, its sine and cosine) and the row-type flag; for the others Table 3 counts splits. A split count measures how often the trees consult a feature, and it favours continuous features with many cut points over binary flags, so it ranks families roughly. Two facts in it are exact. Ten features are never split on and cannot affect the output: the four season flags, the holiday flag, the four feedback aggregates and the family-event flag. And month carries 274 of the 800 first splits, a reflection of a corpus collected in March, April and May, which the retrain answers by dropping the month family.

Training and serving, side by side. Four differences between how v2.6.0 was trained and how model mode would serve it are known and recorded: the row-type flag is always 1 at serving; in October, month takes a value the spring corpus never held, and the climate table covers March to May, so the temperature anomaly is served as 0; the neighbour arithmetic became the training definition on 26 September, after the model shipped; and the shipped model's weekly rows each carried one weather reading from the moment the week was collected. None of the four touches the number production serves today, which reads no model.

5.2 The serving path

For a venue v at hour h, production computes the number in this order.

1. *Fallback.* With no loaded model, or no positive curve value for the venue at that hour, the rule engine answers from the hour, weekday, rating, review count, price level, place types and weather.
2. *The weekly curve* $b(v, d, h)$, rounded and clipped to $[0, 100]$.
3. *The trailing offset*, at weight 0.85: the median of (*reading – curve at the reading's hour*) over the venue's newest 20 live readings taken strictly before h and within 28 days, used with at least two readings and a table rebuilt within seven days, clamped to ±50.

4. *The nowcast.* The newest live reading from an earlier hour, if at most 12 hours old, pulls the number toward itself with weight 1.0, 0.85 or 0.25 at an age of one, two or three hours, and 0 after that.
5. *The level. Quiet* up to 20, *Not Busy* to 39, *Steady* to 69, *Busy* to 84, *Packed* above.
6. *The confidence.* The card publishes the within-15 rate measured for the arithmetic that produced the number, on the 30-day window of Section 6.1: 42.7% for curve and offset alone (its rate over all 119,939 readings), and 83.0%, 40.7% or 38.6% when a reading one, two or three hours old entered. Until 8 October it published the three-day figures (41.7%, 87.4%, 62.8% and 43.5%), and the month brought the three reading-age figures down, the two-hour figure most.

In model mode, step 2 becomes the curve plus the model's departure clamped to ±50, a monotone display map (40 knots, fitted to v2.6.0's output) follows, rounded scores outside 25 to 65 move one point toward the nearer end of the scale, the offset enters at weight 0.5, and the nowcast uses weights fitted on the model's own output. Live readings come from an hourly sweep, at seven minutes past each hour, of active Philadelphia and Lehigh Valley venues, plus Miami venues at two fixed hours a day; curves and offsets are rebuilt after each sweep.

At start-up the server refuses any model whose metadata does not record a passing ship gate, whose feature list names a feature serving cannot compute, whose ONNX input or output shape disagrees with its metadata, or whose published accuracy record is malformed, and it then serves the rule engine.

5.3 What the model never sees

No input identifies or locates a person. The coordinates behind the vector are the venue's, and even those are excluded as features: latitude reaches the model only through daylight, the climate band, the special-night rule and the neighbour box. A signed-in user's account number reaches the predictor only as the key of a per-user rate limit and is never part of the vector, and the training export has no user column. In-app feedback enters as four per-venue aggregates, which the shipped model never splits on. No crowd service reads data from Flux, the sensor of Sections 9 to 18.

6 Results

6.1 H1: the served number on held-out live readings

Table 4. The served number and its comparators on 119,939 live readings over 30 dates (2026-09-09 to 10-08), none of which informed the offset and nowcast weights (fitted on 09-01 to 09-05). Rows (a) and (b) are persistence baselines; (b) was found on this window, so its row is not out of sample. Production and the baselines carry 95% date-block intervals. Bold marks the best value among the configurations scored on the 30-day window. Bias is the mean of predicted minus observed. The last row is production on its first, three-day test, for reference.

Configuration	Within one level ↑	Exact level ↑	Within 10 ↑	Band MAE ↓	MAE ↓	Bias
Production: curve + offset + nowcast	78.4% (77.9–78.9)	**54.9%** (53.9–55.7)	**55.2%** (53.9–56.2)	0.772	18.00 (17.73–18.28)	−1.24
(b) Hour-old reading, else curve + offset	**79.0%** (78.3–79.6)	54.7% (53.6–55.5)	54.9% (53.4–56.0)	**0.756**	**17.68** (17.38–18.03)	−1.16
(a) Reading up to 12 h old, else popular times	74.3% (73.4–75.0)	53.2% (52.0–54.1)	55.1% (53.6–56.2)	0.866	20.05 (19.67–20.50)	+0.48
Curve + offset, no nowcast	70.5%	34.4%	31.2%	1.058	24.55	−1.45
Weekly curve alone (popular times)	64.8%	30.2%	30.9%	1.155	26.40	+4.89
Rule engine	60.6%	22.1%	19.6%	1.287	29.63	+3.95
Before: v2.6.0 as served (map on)	58.3%	30.8%	26.9%	1.337	30.94	+13.94
Constant "Not Busy" (reference)	*75.2%*	*14.8%*	*24.5%*	*1.270*	*28.37*	*−11.82*
Production on the first test (4,183 readings, 09-06 to 09-08)	*78.8%*	*56.0%*	*53.8%*	*0.735*	*17.06*	*−0.98*

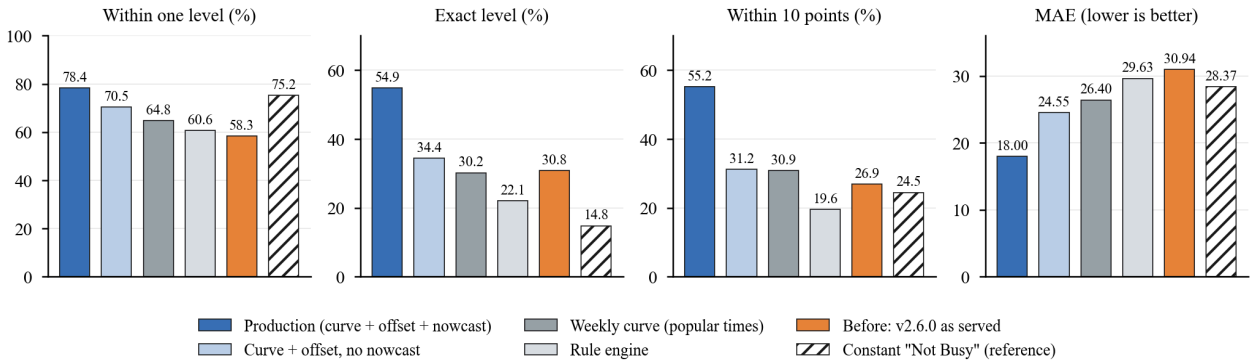


Figure 6. The served number against its comparators on 119,939 live readings over 30 dates (9 September to 8 October 2026).

Table 4 and Figure 6 give the result. Against the configuration it replaced, the served number gained 20.0 points within one level, 24.1 on the exact level and 28.3 within 10 points; MAE fell 42% and the bias fell from +13.9 to −1.2. Against the popular-times baseline, MAE is 31.8% lower and the exact level is 24.7 points higher. The constant row sets the floor. Real venue-hours are bimodal (36.6% *Quiet* and

17.0% *Packed* over the window), and "Not Busy" sits within one level of three of the five levels, so a constant scores 75.2% within one level; it names the exact level 14.8% of the time, with a band MAE of 1.270 against production's 0.772. On the three-day test the constant tied production on within one level; over the month production clears it by 3.1 points (1.1 to 5.1).

Across the month. The three days it was shipped on scored 56.0% on the exact level; the 30 days since score 54.9%. By period, production names the exact level 52.5% of the time from 9 to 25 September (39,000 readings, before the switches went live), 55.7% from 26 September to 2 October (53,287, live in production) and 56.7% from 3 to 8 October (27,652), after the provider began returning fewer live readings (Section 7.2). The rise comes from the mix of readings: the share with a reading from the hour before went from 35.3% to 55.5% and 52.6% as the sweeps sped up. Holding that mix at the month's average, production's exact level is 59.0%, 53.0% and 54.8% in the three periods, lower after 25 September, when the provider's readings became less sticky. Popular times stays between 29.2% and 32.2% on the same readings in all three periods, so production's lead holds in every period. By weekday the exact level runs from 54.4% to 56.7%, Friday and Saturday included, against 28.0% to 32.5% for popular times. Miami is read at two fixed hours a day, so no reading there is ever an hour old and its served number is curve plus offset alone; the nowcast's gain is a Pennsylvania result, measured on 112,728 readings.

Table 5. Production's margin over each comparator on the same 119,939 readings, in points, with 95% date-block bootstrap intervals over the 30 dates. A negative MAE margin favours production. † The interval includes zero under the two-way venue × date bootstrap (exact −0.05 to 0.58; within 10 −0.05 to 0.69).

Production minus	Within one level	Exact level	Within 10	MAE
(b) Hour-old reading, else curve + offset	−0.58 (−0.79 to −0.30)	+0.24 (0.03 to 0.51)†	+0.30 (0.02 to 0.65)†	+0.32 (0.15 to 0.44)
(a) Reading up to 12 h old, else popular times	+4.1 (3.6 to 4.7)	+1.7 (1.3 to 2.2)	+0.1 (−0.1 to 0.4)	−2.05 (−2.31 to −1.85)
Weekly curve alone (popular times)	+13.6 (12.3 to 14.7)	+24.7 (23.1 to 25.8)	+24.3 (22.3 to 25.7)	−8.40 (−9.02 to −7.70)
Curve + offset, no nowcast	+7.9 (6.9 to 8.7)	+20.5 (18.9 to 21.7)	+24.1 (22.0 to 25.5)	−6.54 (−6.96 to −6.05)
Rule engine	+17.8 (16.7 to 18.6)	+32.8 (31.3 to 33.9)	+35.6 (34.0 to 36.8)	−11.63 (−12.08 to −11.02)
Before: v2.6.0 as served	+20.0 (18.7 to 21.5)	+24.1 (22.6 to 25.4)	+28.3 (26.5 to 29.6)	−12.93 (−13.72 to −12.14)
Constant "Not Busy"	+3.1 (1.1 to 5.1)	+40.1 (39.0 to 41.0)	+30.7 (28.9 to 32.0)	−10.37 (−11.33 to −9.45)

Against popular times, curve plus offset, the rule engine, the configuration it replaced and the constant answer, every interval in Table 5 excludes zero, and the two-way bootstrap changes none of these conclusions. That includes the within-one-level margins over popular times and over curve plus offset,

which the three-day test could not settle, and the margin over the constant answer, which it tied. **H1 is supported.**

How much of the gain is persistence. Where the venue was read in the hour before, production serves that reading unchanged, so on those rows it is persistence. Those rows, 48.3% of the window, carry 92% of the exact-level gain over popular times and 96% of the MAE gain. On 24.7% of all readings the value had not changed since the hour before, and the served number is exact on every one of them by construction.

The persistence baselines make the same point from the other side. Rule (a), which knows nothing but the latest reading and popular times, reaches 53.2% exact against production's 54.9%. Rule (b), the hour-old reading or else curve plus offset, matches production: it is 0.6 points better within one level and 0.32 better on MAE, and production's 0.2-point edge on the exact level is inside the two-way interval. The difference between them is the two- and three-hour nowcast weights, which Section 6.3 shows no longer helped over the month. What the curve and offset add is the number for the 52% of readings with no hour-old reading: there, the offset is the only correction, and it lifts curve plus offset over popular times by 5.7 points within one level and 1.85 MAE across the window.

Persistence is not the whole story. On the 59,681 readings whose value differed from the venue's previous reading, production still names the exact level 41.0% of the time against 30.0% for popular times, and on the 28,325 changed readings with a reading from the hour before, 52.9% against 29.1%.

6.2 H2: the delta model on cities it never trained on

Table 6. Point accuracy on the 67,249-row gate slice, scored as the v2.6.0 ship gate scored it: the predicted delta clamped at ±30 and not rounded, with $o_v = 0$ because no held-out venue has earlier readings (display map off).

Predictor	MAE ↓	R^2 ↑	Within 10 ↑
Weekly curve alone	31.48	−0.075	19.2%
Incumbent v2.5.0	30.77	−0.043	19.3%
Delta model v2.6.0	**29.42**	**+0.040**	**20.7%**

The model improves on the curve by 2.06 MAE (6.5%) and 0.11 R^2, and on the incumbent by 1.35 MAE. R^2 crosses zero: this is the first version to predict these cities better than the mean of the labels. The gain holds in every held-out city (Figure 7): 2.26 MAE in Miami, 2.21 in Barcelona and 1.72 in Tokyo, where within-10 is 23.6%, 20.9% and 16.9%. Every number in this section reproduces from the saved model and holdout (the incumbent from its tagged commit).

The same three cities have been the holdout since March and scored every version from v2.2 on, and the score chain's clamp and one-point push were fitted on the first 30% of their dates. They therefore serve as a validation set: the margins in this section belong to the version that passed after earlier versions failed, not to one untouched test. The October retrain is scored on a time holdout instead (Section 7.2).

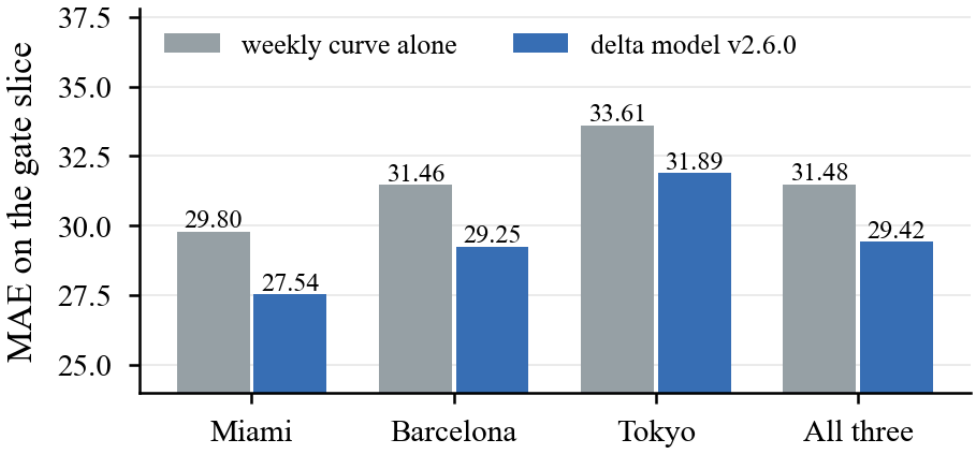


Figure 7. Delta model against the weekly curve on each held-out city of the gate slice.

On the same rows, measured on the five-level scale (Table 7), the model wins on every column: +2.6 points of exact level, +4.3 within one level, and 0.086 fewer levels of error on average.

Table 7. Level accuracy on the gate slice, with the delta clamped at ±50 and not rounded, and a score of exactly 20 counted as *Not Busy* (display map off). Production counts 20 as *Quiet*; on that ladder the gains are +2.5 exact, +4.4 within one level and 0.089 fewer levels.

Predictor	Exact ↑	Within one ↑	Band MAE ↓	MAE ↓
Weekly curve	22.9%	57.7%	1.315	31.48
Delta model	**25.5%**	**62.0%**	**1.230**	**29.35**

Scored exactly as Equation (3) writes it, with the ±50 clamp and rounding, the model reaches MAE 29.35, R^2 +0.042 and within-10 21.7%. Rounding to whole numbers, against labels on a 5-point grid, adds a point of within-10: the grid effect of Section 3.3 again.

Table 8. The delta model's margin over the weekly curve on the gate slice, with 95% intervals from a paired two-way cluster bootstrap (2,000 resamples, stratified by city) that resamples venues and city × month × weekday blocks. Each metric is scored on the chain of its own table.

Held-out rows	MAE	Within 10	Exact level	Within one level
All three cities	−2.06 (−2.75 to −1.44)	+1.51 (0.44 to 2.65)	+2.61 (1.55 to 3.79)	+4.31 (2.93 to 5.80)
Miami	−2.26 (−3.50 to −1.13)	+1.87 (−0.01 to 3.93)	+2.62 (0.65 to 4.78)	+4.83 (2.39 to 7.70)
Barcelona	−2.21 (−3.22 to −1.27)	+2.54 (0.38 to 4.87)	+4.27 (2.19 to 6.55)	+3.68 (1.79 to 5.76)
Tokyo	−1.72 (−2.69 to −0.83)	+0.50 (−0.66 to 1.85)	+1.71 (0.46 to 3.23)	+3.98 (2.01 to 6.17)

The holdout carries no venue or date column, so both clusters had to be recovered. Venues come from their static attributes: the daylight feature is an exact function of latitude and month, which identifies 2,644 venues, each holding exactly one row per weekly slot. City × month × weekday is the finest date proxy the rows carry. The dependence is temporal: date blocks widen the MAE interval 7.4 times over

treating rows as independent, where venues alone widen it 1.4 times. R^2 gains 0.115 (0.079 to 0.154), band MAE falls 0.086 (0.058 to 0.116), and against the incumbent v2.5.0 MAE falls 1.35 (1.00 to 1.72). Every margin excludes zero overall and in every city except within-10 in Miami and in Tokyo taken alone, so the claim for H2 rests on MAE and the level metrics.

Out of fold, where every training row is scored by a model that never saw its city, the 369,076 served realtime rows give MAE 27.54, R^2 0.127, within-10 22.8% and within-15 33.3%. Refit with the same recipe and scored 15 days forward within the 30 training cities (103,254 realtime rows), the model scores MAE 29.02 against 30.23 for the curve (within-10 20.7% against 20.4%), and R^2 0.078 against 0.016. The gain is smallest in the dinner (17:00 to 20:00) and evening (21:00 to 23:00) hours, 0.96 and 0.87 MAE.

Placed against prior work, the 6.5% comes with no recent history, where BysGNN's larger reductions read the previous 24 hours at every venue [7]. Its R^2 of +0.040 is small, at a granularity where the nearest peer-reviewed occupancy study's section-level forecasts a day or more ahead have R^2 near zero and often below it [2]. **H2 is supported on the held-out cities**, with the selection caveat above. The model learned from spring rows of unknown provenance, and on verified live readings the online correction of Section 6.1 gives the stronger number, which is why production serves it (Sections 6.1 and 19.5).

6.3 H3: the value of a reading by its age

The served number's accuracy is set largely by how recently the venue was read (Figure 8). Over the 30-day window, 48.3% of readings (57,900 of 119,939) had a reading from the hour before, and on those the served number lands within 15 points 83.0% of the time, within one level 88.6%, names the exact level 77.0% of the time and has an MAE of 9.99. At two hours within-15 falls to 40.7% (14,778 readings) and at three to 38.6% (7,794), level with curve and offset on the same rows (39.7% and 38.7%); with no reading in the last 12 hours it is 44.4%. On the three-day test the two-hour figure was 62.8%: over a month nearly all of a reading's value sits in its first hour. The card had published that 62.8% as its confidence for a two-hour reading, and on 8 October it was replaced with the month's figures (Section 5.2, step 6).

On other metrics the two-hour weight of 0.85 did harm. On those 14,778 readings production's MAE is 28.56 against 25.83 for curve plus offset alone (2.73 worse, interval 1.41 to 3.63) and its within-one rate 4.5 points lower (−6.1 to −2.4); it is no better than popular times on MAE (27.34). The weight was fitted on 1 to 5 September, when consecutive readings were identical four times in five. At three hours the 0.25 weight makes no measurable difference. **H3 is supported:** the decay in accuracy follows the decay in correlation measured in Section 3.2, and is steeper. Production has not been changed before the October gate, which compares a retrain against production as served; rule (b) is reported beside that verdict (Section 7.2).

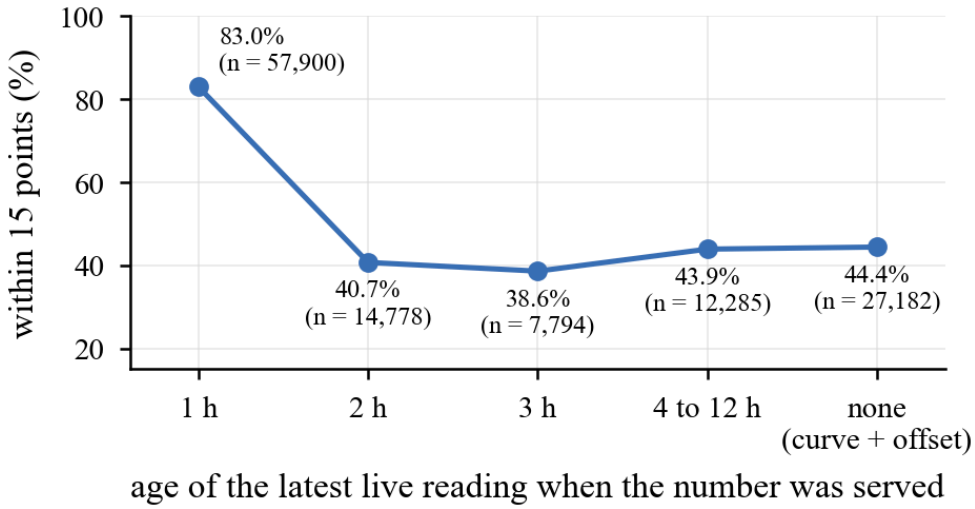


Figure 8. Accuracy by the age of the latest reading, over the 30-day window (119,939 readings). The card publishes the one-, two- and three-hour rates as its confidence and, when no reading under four hours old entered, 42.7%: the within-15 rate of curve and offset alone over all 119,939 readings (Section 5.2).

6.4 What coverage is worth

A delta model cannot score a venue that has no curve, so I compared the model path with the rule engine that serves such venues. Both were given identical inputs on 5,025 verified live readings in Pennsylvania at venues that hold their own curve, with the events lookup disabled for both (Table 9, Figure 9).

Table 9. The two serving paths on the same 5,025 live readings (display map on, as served then).

Path	MAE ↓	Within 10 ↑	Within 15 ↑	Median error ↓	On-grid
Model path	28.22	**31.7%**	**41.5%**	21	41.5%
Rule engine	30.04	20.3%	29.2%	27	19.8%
Stored curve alone	**26.61**	30.0%	**41.5%**	**20**	100%

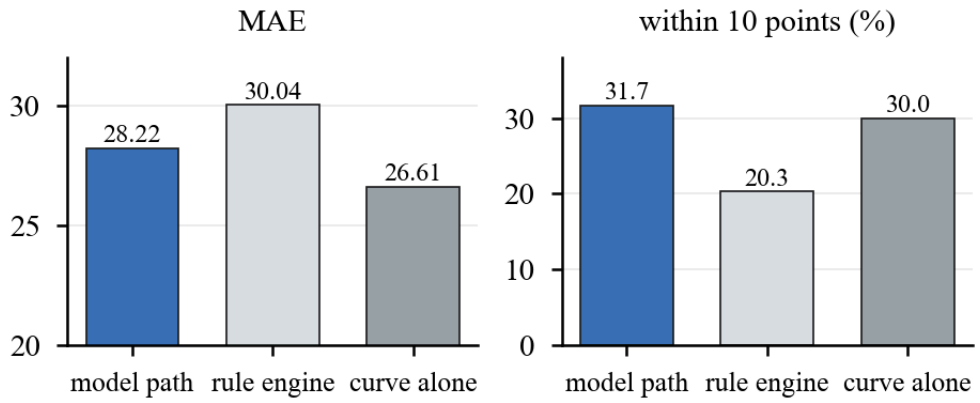


Figure 9. The model path and the rule engine on the same 5,025 live readings, with the curve alone for reference.

Coverage is worth 1.82 MAE, 11.4 points of within-10 and six points of median error. Both arms ran with the events lookup disabled, so the comparison leaves the value of events unmeasured. Grid-equalised, coverage is still worth 11.2 points of within-10, so the result is not a lattice effect. The bare curve beats the model path on MAE by 1.61 while losing on within-10: on these bimodal labels the two metrics reward different predictions (Section 7.1).

6.5 Ablation experiments

Nowcast and offset. Table 4 separates the two corrections:

- *Smoothing and the offset* together take the popular-times row from 64.8% to 70.5% within one level, from 30.2% to 34.4% exact and from 26.40 to 24.55 MAE, and move the bias from +4.89 to −1.45. Within-10 ends where it began (30.9% to 31.2%): smoothing moves predictions off the 5-point grid and lowers it to 26.4%, and the offset lifts it back, the grid penalty of Section 3.3 in miniature.
- *The nowcast* then does most of the work: within-10 rises from 31.2% to 55.2%, and the exact level from 34.4% to 54.9%.

Rescaling the curve directly. The obvious remedy for the poor calibration of Section 3.1 is to rescale the curve toward the mean. I priced it on the 9,903 paired rows (Table 10, Figure 10).

Table 10. Corrections of the curve's calibration on 9,903 paired rows. The two per-venue rows are fitted and scored on the same readings, so they are in-sample estimates, not held-out results.

Predictor	MAE ↓	Within 10 ↑
Curve alone	26.86	31.6%
Curve + one global mean offset	25.88	28.9%
Global slope correction	25.49	21.8%
Per-venue intercept, half weight (in-sample)	23.18	**35.9%**
Slope + per-venue intercept (in-sample)	**22.55**	27.2%

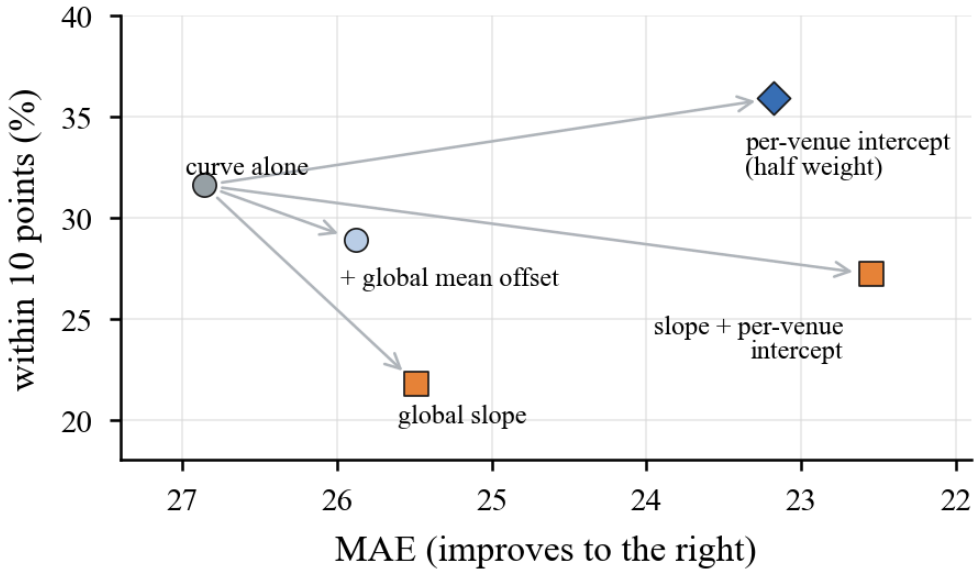


Figure 10. Pricing the corrections of Table 10. Corrections that narrow the spread (squares) buy MAE and spend within-10. The per-venue intercept, the form of Equation (1), here at half weight, improves both.

Every correction that narrows the published spread buys MAE and spends within-10, so the system uses the per-venue intercept. A leave-one-out test over 688 venues (11,804 held-out readings) repeats the pattern out of sample: a per-venue slope buys 0.99 MAE over the per-venue intercept and costs 3.19 points of within-10, while the intercept itself buys 1.56 MAE against the bare curve (25.97 to 24.41) at a cost of 0.1 point of within-10. On 8,895 September live readings, the half-weight offset moves MAE from 25.90 to 24.53 and the exact level from 32.1% to 38.0%, on five of five days and in both cities.

The display map. From 2026-08-28 to 09-26 a 40-knot monotone quantile map sat between rounding and the level cut. On the 46,101 later held-out-city rows it was not fitted on, it bought 8.4 points of within-10 for 3.15 MAE. September's live readings reversed that verdict: the map cost 10.5 points of within one level (57.4% with it against 67.9% without, on 7,920 readings), and on 26 September it left the served path, because curve plus offset bypasses it (model mode would still apply it).

Volume against depth. A paired test scored two models on the same 17,322 later rows. One was trained with the scored venues' earlier history, 304,000 more rows in total; the other had never seen those venues. The model with the extra history scored 27.63 MAE against 27.14, worse by 0.49, in a single paired run without an interval, so it is a pointer rather than a measurement. A two-head variant makes the same point from the other side:

- trained on the unknown-provenance pool, it beats the shipped model by 0.88 MAE (95% interval 0.51 to 1.27);
- held to the 4,223 verified labels available at the time, it loses by 1.59 MAE (interval 0.86 to 2.45).

The limit is the supply of verified labels per venue, and the live readings of Section 6.1 are the input that moves it.

The learner. To test the choice of XGBoost I held the features, the target and the score chain fixed and changed the learner alone, in a rotation across the three held-out cities: each learner trained on two and was scored on the third (Table 11). Each fold gives a learner 135,000 to 165,000 rows from two cities, far fewer than the shipped model's 1.93 million from 30, so the shipped model appears for reference. Neither boosted model was tuned, seeds and thread counts were pinned, and two runs gave identical predictions.

Table 11. Four learners and the weekly curve, with the same 106 features, target and score chain (Equation (3)), each learner trained on two held-out cities and scored on the third, pooled over the 67,249 gate rows. Bold marks the best learner in each column.

Learner	MAE ↓	R^2 ↑	Within 10 ↑	Exact ↑	Within one ↑
XGBoost, the settings of Section 4.4	29.97	0.000	21.4%	25.1%	60.7%
LightGBM	**29.80**	**+0.010**	**21.7%**	**25.3%**	**61.1%**
Ridge regression	30.18	+0.003	20.3%	24.1%	59.9%
Category-by-hour mean delta	31.61	−0.095	19.7%	23.3%	57.7%
Weekly curve alone	31.48	−0.075	19.2%	22.9%	57.7%
Shipped v2.6.0 (30 other cities)	*29.35*	*+0.042*	*21.7%*	*25.4%*	*62.0%*

Against XGBoost, LightGBM is better by 0.18 MAE (0.07 to 0.29) and level with it on within-10; ridge regression is within 0.21 MAE (−0.02 to 0.49) but loses 1.09 points of within-10 (0.54 to 1.72) and 1.01 of the exact level; a category-by-hour average is no better than the curve. Among the learners tested, the gradient-boosted ensembles did best: LightGBM had the lower MAE and XGBoost matched it on within-10. The shipped model beats every learner here by at least 0.45 MAE, but it trained on far more rows from 30 cities against two, so the comparison does not separate breadth from size.

6.6 The retrain, rehearsed

Scored on the first test's 4,183 readings (6 to 8 September), a model trained with five live features on 3,579 earlier live readings reached 79.7% within one level, +21.2 points over v2.6.0 as served (95% interval 15.8 to 24.2). The features were the last reading's deviation and age, the trailing offset and its count, and the curve one hour earlier; they carry 47% of its split gain and are worth +10.0 points of within-10 over the same candidate without them. It trailed production by 12.2 points of within-10 (interval −16.3 to −2.7) and by 1.9 MAE. Treating the 4,183 readings as independent samples instead gives intervals of 19.3 to 23.1 and −14.3 to −10.1; the date-block intervals are wider by factors whose squares are about 5 and 10, the first-test clustering estimates that Section 7.2 uses for its power arithmetic. The band gate refused it on the MAE guard and the three-date sample, and the within-10 criterion of Section 4.4 was added so that a loss like this one is refused on its own.

7 Discussion

7.1 Interpretation

The latest reading is the lever. Adding the venue's latest live reading raises within-10 from 31.2% to 55.2% and the exact level from 34.4% to 54.9% over 30 days (Section 6.5); with the offset, the served number cuts the curve's MAE by 31.8%. Most of that is persistence. Where an hour-old reading exists, repeating it is the best this system found, and the persistence baselines come within two points of production on the exact level (Section 6.1). The curve and its offset matter where no fresh reading exists, which is half the time. For anyone building on typical-week curves, persistence of the latest reading is the baseline to beat, as it is in forecasting generally [8], and the curve is the fallback.

The model's gain holds in every held-out city. A 6.5% MAE reduction on cities that also served as its validation set (Section 6.2), with R^2 crossing zero at single-venue, single-hour granularity, is what the spring data supports without any observation of the venue's recent state. On the gate rows the model's output correlates with the raw curve at 0.9638: the curve supplies the rhythm and most of the level, and the model supplies a correction.

Evidence density is the binding constraint. The spring corpus holds about 25 realtime rows per venue across 168 weekly slots (26.3 after the September merge). On 2026-09-06, 9,339 rows across 1,228 venues carried verified live provenance, a median of six each. 304,000 more rows of the same kind did not help in the single paired test of Section 6.5, which has no interval; the hourly live collection is the depth the problem asks for, and a reading's value lasts about an hour, so how often a venue is read matters more than how long its history is.

One number cannot sit near both modes. Over the 30-day window, 36.6% of readings are *Quiet* and 17.0% *Packed*: a bimodal target on which MAE and within-10 can give opposite instructions, as the curve corrections of Section 6.5 show. A live reading moved both at once: adding the nowcast to curve plus offset cut MAE from 24.55 to 18.00 and raised within-10 from 31.2% to 55.2%.

Who it helps. For a group choosing where to go, moving the exact level from 30.8% to 54.9% turns the card into an answer that is right more often than not. For a venue, a number people trust about the present can make a quiet hour worth visiting and move some demand off the peak. For anyone building on typical-week curves, three findings carry over without any of this code: a typical-week curve used as a forecast is poorly calibrated and should be judged as a forecast and against persistence of the latest reading; a within-k metric on lattice-valued labels should be read beside its on-grid rate; and fresher readings of each venue help more than more rows of the same kind. Flock receives only venue-level scores from the busyness provider, which builds them from aggregated, anonymised signals; no individual's data reaches Flock.

7.2 The October retrain

This paper is a snapshot taken before the mid-October retrain. That retrain will learn the live features rehearsed in Section 6.6 rather than blend a reading in after the fact, train with live readings carrying 80% of the loss, drop the month family and the row-type flag, and add the sports-schedule features. It will be gated against production as served at gate time, curve plus offset plus nowcast, and shipped only if it wins; rule (b) of Section 6.1 and the constant answer will be reported beside the verdict. The

run is set for about 19 to 20 October, and I will update this paper with its held-out results. That update adds the strongest test in this project: a model trained on a month of verified live readings, scored on a 14-day window whose rules were fixed before its data existed.

The sample. On 25 September, at about 2,100 readings a day, I projected about 90,000 to 95,000 provenance-verified live readings by mid-October, with the last 14 days held out. From 25 September the sweeps ran in parallel, and from 26 September to 2 October the window of Section 6.1 gathered about 7,600 model-served readings a day. From about 23:00 on 3 October the provider answered the same hourly calls with fewer live readings, about 4,600 a day from 3 to 8 October, 61% of the week before, and since about 16:00 on 12 October it has returned none; the provider has been asked why. The registered rule anchors the 14-day window on the last date with live readings. If the feed does not return before the run, the window is 29 September to 12 October, which holds far more than the floor of 1,000 readings over 5 dates. Ten of those dates fall inside the month reported in this paper, so production's accuracy on them is already known. The retrain trains only on earlier dates, and none of its settings were chosen on these ten days. Among the readings the provider did return after 3 October, the drop did not lower accuracy: the share with a reading from the hour before moved from 55.5% to 52.6%, and production named the exact level 56.7% of the time after the drop against 55.7% before, or 54.8% against 53.0% with the mix of reading ages held at the month's average (Section 6.1). Readings that were never returned cannot be scored. The decision rule does not change. The 30-day test of Section 6.1 is the last measurement of the current method before the retrain, and the retrain is gated against that method as it serves.

A pre-registered test. Hypotheses H1 to H3 were stated after the analyses that test them (Appendix A). The October retrain is the first test in this project whose hypotheses, data split and decision rule were fixed before its data existed. They are fixed in code: the band gate's thresholds and the 14-day time holdout were committed on 2026-09-25 and 2026-09-26, three weeks before the held-out window closes.

- **H4.** The retrain scores higher than production as served on within one level over the held-out 14 days, and its date-block 95% lower bound is above −1.0 point.
- **H5.** It is not worse on within 10, by the same bound.
- **H6.** It beats the weekly curve on within one level with a lower bound above zero, and its band MAE is no worse than the curve's.
- **Guards.** At least 1,000 readings over at least 5 dates; production never saw the window; no city with 300 or more readings falls more than 2 points behind production; within one level is not below the rule engine's; MAE rises by at most 1 point.
- **Analysis.** 2,000 date-block bootstrap resamples. The constant-answer reference and persistence rule (b) are recorded beside every verdict, as references only; neither changes the decision.
- **Decision.** The retrain ships only if all of these hold. Otherwise production keeps serving, and the result is reported either way.

Table 12. The smallest difference the held-out window can detect at 80% power (two-sided, 5%), in points, by the number of held-out readings. Each cell gives the value with readings treated as independent, then with the variance inflated 5 and 10 times, the clustering the first-test rehearsal of Section 6.6 measured, used here as an illustration; the 30-day window measured about 15 times on the exact-level margin over popular times (Section 4.5). The assumed rates are production's over the 30-day window. The last column gives the readings needed to detect a 1-point difference with independent readings.

Metric (assumed rate)	15,000 held out	29,000 held out	60,000 held out	Readings for 1 point
Within one level (78.4%)	1.33 / 2.98 / 4.21	0.96 / 2.14 / 3.03	0.67 / 1.49 / 2.10	26,600
Exact level (54.9%)	1.61 / 3.60 / 5.09	1.16 / 2.59 / 3.66	0.80 / 1.80 / 2.55	38,900
Within 10 (55.2%)	1.61 / 3.60 / 5.09	1.16 / 2.59 / 3.66	0.80 / 1.80 / 2.54	38,800

The arithmetic of Table 12 shows how strict the gate is. With about 60,000 held-out readings, an illustrative size, a 95% interval on a within-10 difference is about ±0.6 points if readings were independent, and ±1.3 to ±1.8 points with 5 to 10 times inflation. At that inflation a retrain that exactly ties production would see its lower bound fall below −1 point and be refused; to clear it, it would need a within-10 win of about 0.3 to 0.8 points. The gate scores paired date-block differences, whose interval depends on how the two systems disagree across readings and dates, so the real requirement may be lower. I have kept the bound strict on purpose: the rehearsal in Section 6.6 is the case it exists to stop.

8 Conclusions on the forecast

The number Flock serves today corrects the venue's curve with its recent live offset and a nowcast of its latest reading. On 119,939 live readings its weights never saw (9 September to 8 October 2026) it named the exact crowd level 54.9% of the time, landed within one level 78.4% of the time (a constant answer reaches 75.2% but is exact only 14.8% of the time) and came within 10 points 55.2% of the time; the configuration it replaced scored 30.8%, 58.3% and 26.9%, and the served number's MAE is 31.8% below the popular-times baseline's. The three-day result it was shipped on held up on 29 times as many readings, across all seven weekdays and every period measured. Most of the gain is persistence: an hour-old reading puts 77.0% of served numbers on the exact level, and a rule that repeats it, or else serves the corrected curve, matches production. The freshest reading of the venue is worth more than any correction to its curve. Separately, a gradient-boosted delta model, which production does not serve today, beats the curve on three held-out cities, which also served as its validation set, without any recent-history input.

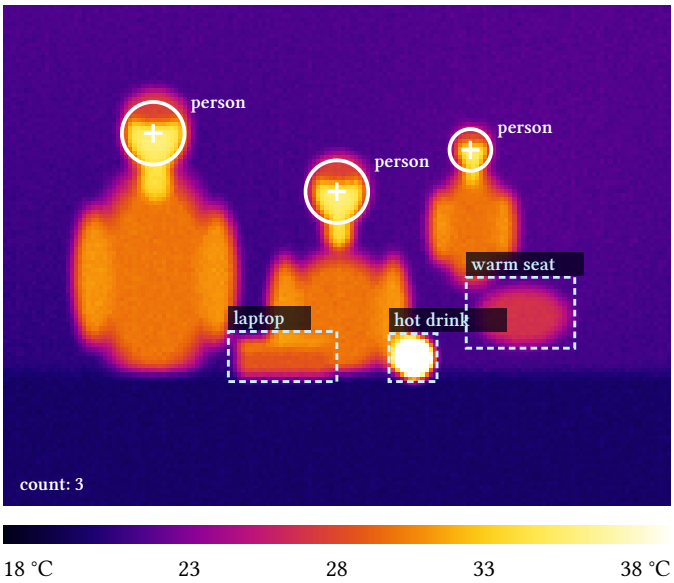
Three findings should carry over to other systems built on typical-week curves, each measured here on one provider's readings:

- as forecasts, those curves are poorly calibrated, with the observed reading rising about half a point per point of curve, so they are judged against persistence of the latest reading;
- within-k metrics on lattice-valued labels reward predictions that land on the lattice, so they are read beside the on-grid rate;

- the freshest reading of the venue carried most of the gain over popular times, and its value was spent within about an hour.

9 Counting people on a thermal camera

Flux is a small venue sensor, built around a Raspberry Pi 5 and a FLIR Lepton 3.5 thermal camera, that a venue could install to count the people actually there; none is installed at a venue yet, and Flock's crowd number does not read sensor data. A thermal camera records heat instead of light. At 160 × 120 thermal pixels Owl is built to count people, never to identify them, and Flux reduces each frame to a count on the device and throws the frame away (Figure 11).



Camera	FLIR Lepton 3.5, 160 × 120, 57°
Input	degrees Celsius, every pixel
Model	718,550 parameters
On the sensor	14–16 ms a frame, 1 Pi 5 core
Names	person + 9 kinds of warm thing
Counts	people only
Leaves the device	counts and a sound level
Footage I recorded	none

Figure 11. What Owl reads and what it reports. An illustrative 160 × 120 frame, drawn for this paper the way Owl's training generator draws a scene, shown at the camera's true resolution in the iron palette the Lepton uses. Owl puts a point on each person's head and a named box on everything else warm. Only the points are counted: this frame counts as three.

The first counter on Flux was a rule. It found regions warmer than the room, merged them on a coarse grid, and counted the ones large enough to be a person. The rule failed in predictable ways: a hand held up to the lens counted as a person, spread fingers counted as two, a mug of coffee or a laptop could count as somebody, and a packed room merged into one warm mass. Telling these apart is a question of shape, which is what a learned model does well.

Sections 9 to 18 describe Owl, the model that replaced the rule, and what building it showed. Four constraints shaped the work from the first day.

1. **It runs on the sensor.** One core of a Raspberry Pi 5, with no GPU and no round trip to a server.
2. **It names kinds of things and nothing more.** No identity, no face features, no tracking from one frame to the next. Only numbers leave the device: people counts and a sound level.
3. **I do not record people to train it.** Owl learns from frames I generate and from existing datasets.

4. **Every real frame is licensed.** Each source allows commercial use under its licence, or its authors gave written permission.

My contributions are a generator of radiometric thermal scenes matched to the Lepton 3.5 (Section 12.2); a 718,550-parameter network that finds people as points and names nine other kinds of warm object, decoded on the sensor with physical sanity checks (Sections 12.4 to 12.6); a crowd second look using a density head whose total does not shrink when heads touch (Section 12.6); an evaluation on six real benchmark sets with block-bootstrap intervals, a 41-case stress test and a per-situation blind-spot report, with the original rule scored on identical frames, as deployed and tuned to each set (Section 13); and 15 findings from 19 full training runs plus five short trials (Section 15).

10 Background on thermal people counting

Thermal people counting is usually studied with low-resolution array sensors on the ceiling, such as the 32 × 24 MLX90640 in PUT Thermo Presence [25], whose counting models have been benchmarked on resource-constrained hardware [34], and with higher-resolution cameras for segmentation and pose [24, 26]. Owl borrows from three lines of work.

Objects as points. CenterNet [17] predicts a heat map with a peak at each object's centre and decodes by finding local maxima. Counting then reduces to counting peaks, which is exactly the quantity a venue needs. Its penalty-reduced focal loss, introduced by CornerNet [32] as a variant of focal loss [18], handles the imbalance between a handful of peak pixels and thousands of background pixels.

Counting by density. Density-based counting [19, 20, 33] regresses a map whose sum is the count. It degrades gracefully when objects overlap, where peak finding merges two touching heads into one.

Learning from generated data. A model trained on generated frames transfers to real ones only as far as the generator covers reality. Domain randomization [21] varies everything the model should ignore so that the real world looks like one more variation. Section 15 shows where that held for Owl and where it did not.

11 The sensor and the baseline

Flux reads the Lepton 3.5 through a PureThermal 3 board over USB. The camera reports radiometric temperatures in degrees Celsius for 160 × 120 pixels across a 57° horizontal field of view [29]. Through that lens a head three metres away is about 8 pixels wide and a hand 30 cm from the lens about 45.

The baseline rule bins the frame 4 × 4, estimates the room as the 20th percentile of the frame, learns fixed warm spots such as radiators per cell over time, keeps cells warm enough to be skin, groups them by eight-connectivity, and counts groups of at least 6 cells, about one head at three metres. Its settings were tuned on a bench with a real person and against constructed scenes, and it works for well-separated people at moderate range. It has no notion of shape, which is the gap Owl fills. Throughout this paper the rule is scored on exactly the same frames as Owl.

12 Owl: method

Figure 12 shows the pipeline, from the training data to the counter on the sensor.

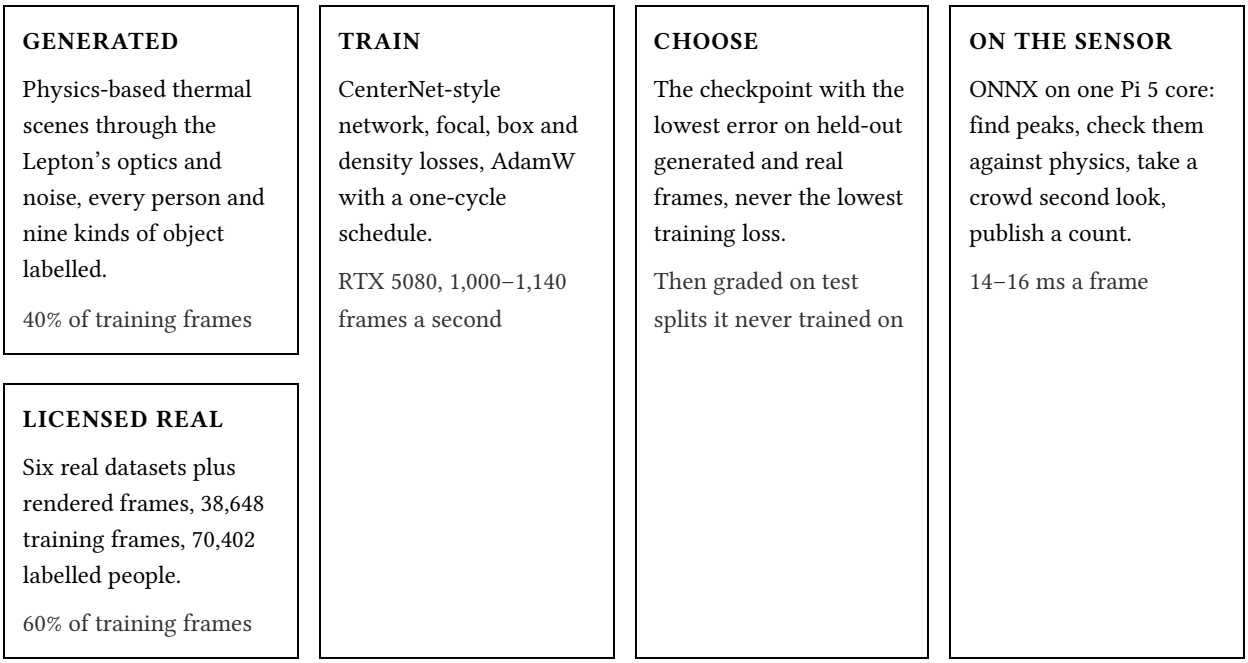


Figure 12. The pipeline. Data on the left, the deployed counter on the right. Checkpoints are selected on validation splits; versions, data shares, decoding settings and the crowd model's hand-over point were kept or rejected on the real test splits, which makes them development benchmarks (Section 13). Accuracy is measured on those splits and on generated stress scenes, and speed on the sensor itself.

12.1 Formulation

Each frame is labelled with one point per visible person, at the centre of the head, or at the centre of the visible body when the head is out of view. Every other warm object gets a box and one of nine kinds: hand, pet, hot drink, food, laptop, screen, heater, lamp and warm seat. Counting people is finding peaks in the person map. The other nine maps let the screen show why something was not counted, and they give the network somewhere to put the things it must not count.

12.2 A generator matched to the camera

Owl's first version learned only from generated frames, and the generator stayed in every training batch after that (40% of frames for the shipped model). It draws a room at some temperature, people at every distance and pose, and the objects that fool a threshold, at twice the camera's resolution, then averages down, blurs for the lens, and adds the camera's noise. Table 13 lists the physical model, which replaced hand-picked temperatures after the first round of real-footage tests (Finding 4).

Table 13. The physical model behind the generated frames. Values were set from published thermography of skin and clothing.

Surface	Model
Skin	33.4 °C forehead in a 22 °C room, plus 0.14 °C for every degree of room temperature
Clothing	A garment-specific share of the way from the room to the skin; a winter coat reads barely above the room
Hair	3–10 °C above the room
Pets	Coat near room temperature, eyes and ears 35–38 °C
A seat just left	4–8 °C above the room
Rooms	Typical rooms, plus 20% of frames drawn hot at 29–34 °C
Camera	Lens blur, noise, fixed-pattern column stripes, dead pixels, vignetting, several degrees of absolute error, gain error, a crooked mount, the frame after a flat-field correction

Scenes include standing, sitting and lying people, children, hats, hoods, glasses, masks, bald heads, bare torsos, winter coats, people hidden behind tables and bars, one person partly behind another, and crowds. They also include things that must not be counted: hands at the lens, reflections of people in glass and shiny floors, warm footprints, sun on the floor, candles, cold drinks and bulbs hotter than the camera's range.

12.3 Licensed real data

Six real thermal datasets and one set of rendered frames entered training over the project (Table 14). Two are from the Lepton 3.5 itself. Sources that store temperatures (PUT, SDT) are used in degrees Celsius. Sources stored as 8-bit images carry no temperatures, so each frame is mapped onto a plausible room-to-skin range drawn afresh every time it is used (low end 16–28 °C, span 8–16 °C, gamma 0.8–1.25), which stops the network from treating any single mapping as truth. IISER's frames are colour images in the camera's iron palette; I convert them back to one channel by luminance, which rises monotonically along that palette from black through purple, red, orange and yellow to white.

Table 14. Real training data. People are counted in the training split; for IISER they are labelled faces. Licences and permissions are listed in the acknowledgements.

Source	Camera and view	Stored as	Train frames	People	Use
PUT Thermo Presence [25]	MLX90640 32 × 24, ceiling	°C	8,472	24,106	MIT licence
OpenThermalPose2 [26]	Side view, indoor and outdoor	8-bit	8,846	15,849	MIT licence
AAU VAP Trimodal [24]	Wall-mounted, meeting rooms	8-bit	3,964	5,466	CC BY 4.0
SenSys 2021 set [27]	32 × 24 rooms with hot laptops	8-bit	1,239	1,577	CC BY 4.0
SDT, TU Wien [22]	**Lepton 3.5**, elevated indoor, one person or none	°C	4,800	3,600	Authors' permission
SDT rendered frames [22]	Rendered to the Lepton's look	°C	8,000	6,000	Authors' permission
Thermal Vision, IISER Bhopal [23]	**Lepton 3.5**, classroom of 12 to 15 and lab of up to 6 (up to 10 and 5 labelled faces)	Iron palette	3,327	13,804	Authors' licence and permission

Real datasets label people and nothing else, so for real frames the loss on the other nine maps is switched off. AAU, SenSys and SDT mark bodies rather than heads; for them the head point is estimated from the box and its target is spread over the pixels the head could occupy (σ = 5 px), a fix that Finding 7 explains. Test splits are the datasets' own, or, where a dataset has none, contiguous ranges in time: SDT by frame range 60/15/25, and IISER by time within each room 60/15/25 with ten frames dropped at every boundary, because neighbouring frames are near copies.

12.4 Network

The network sees two channels. The first is absolute, (t − 30) / 8, so a 60 °C mug can never look like skin. The second is relative to the frame's median, (t − median) / 4, so the several degrees an uncalibrated Lepton reads high or low do not move a person out of range. An encoder halves the resolution four times, to strides 2, 4, 8 and 16 with 24, 48, 72 and 96 channels, using residual blocks [30] with dilation at the deepest stage [31, 33] to widen the view: a hand filling half the frame must be seen as one hand and not five warm blobs. A top-down path fuses the deep features back to stride 4, a 40 × 30 grid. Three heads read that grid: ten heat maps (bias initialised to −4.6, so training starts from "nothing here"), a box head that predicts each object's reach in log space, and for the crowd model a density head through a softplus, initialised near empty. The main model has 718,550 parameters and exports to ONNX as a 2.9 MB file.

12.5 Training

The loss is CenterNet's penalty-reduced focal loss on the heat maps [17, 18], with $\alpha = 2$ and $\beta = 4$ as in CornerNet [32], masked by which maps a frame actually labels; an L1 loss on the log of the box reach, at object peaks; and, for the crowd model, squared error per density cell plus a tenth of the error in the total. I train with AdamW (weight decay 10^{-4}), a one-cycle learning-rate schedule with 5% warm-up, gradient clipping at 5, and batches of 64. Fresh runs use a peak rate of 2×10^{-3}; later versions start from the previous version, keeping every layer whose shape matches, at 2×10^{-4}. Eight CPU workers generate frames on the fly for an RTX 5080. The project ran 19 full training runs plus five short trials, 581,700 steps in all, which is 37.2 million training frames. Each run keeps the checkpoint with the lowest combined error on held-out generated and real validation frames.

12.6 Decoding on the sensor

The sensor runs the ONNX file with ONNX Runtime [13] in full precision on one core. A person is a cell at least as confident as all eight neighbours and above 0.4. Rarer kinds must clear a higher bar before they are named, from 1.25 times the person threshold for hands and hot drinks to 1.75 times for pets, because they are learned from fewer examples. A person peak is dropped when a screen, laptop, heater, lamp or warm seat is more confident in the same 3×3 neighbourhood. Two physical checks run on a 9×9 patch around each person: its warmest pixel must be at least 1 °C above the room, taken as the frame's 10th percentile (Finding 3), and its median must not exceed 45 °C, which no person reaches. When five or more people are found, a second crowd model reads the same frame, and the published count is the highest of the main model's peaks, the crowd model's peaks and the rounded density total.

13 Owl: evaluation protocol

Validation chooses the checkpoint; the test splits chose everything else. Each run keeps the checkpoint with the lowest combined error on held-out generated and real validation frames, never the lowest training loss. Validation frames are resampled with each run's own source weights, so validation scores are not comparable across versions. Every version was then graded on the six test splits, and versions, data shares, decoding settings and the crowd hand-over point were kept or rejected on those grades. The six sets are therefore development benchmarks, and the figures in Section 14 are development results. No test frame was ever trained on: the training cache holds training and validation frames only, with no exact duplicates of a test frame. A sealed test is set out under What comes next.

Six real benchmark sets. Six test splits that no model trained on: the Lepton 3.5, Flux's camera model, in indoor rooms (SDT, 2,000 frames) and in a crowded classroom and lab (IISER, 1,375 frames, 4,882 labelled people), a ceiling camera (2,347 frames, 8,043 people), rooms with hot laptops (265 frames), wall-mounted meeting rooms (813 frames, 1,061 people) and a side view (2,151 frames, 3,872 people). A frame is *exact* when the count equals the truth and *within one* when it is off by at most one person. The rule is scored on the same frames, as deployed and, separately, with its settings tuned on each set's validation split. Intervals are block bootstraps that resample runs of consecutive frames within a scene or sequence (2,000 resamples), because neighbouring frames are near copies.

Two properties of the sets limit what they show. IISER's labels mark visible faces, so a person facing away from the camera is absent from its truth. And some splits share scenes: AAU's three meeting rooms and their people appear in training and test, and one PUT test sequence continues a training sequence. Dropping that sequence and every test frame nearly identical to a training frame raises Owl's PUT score from 92.9% to 96.0% exact and moves no other set by more than 0.1 points, so near-copies do not inflate the results.

Stress test. 41 hard situations, each generated on purpose 150 times from a seed fixed per case and scored alone, counted exactly as the sensor counts, crowd model included. The cases come from the same generator as 40% of the shipped model's training data, and training was steered toward the cases it failed, so the stress test is a tuned, in-distribution check of known hard cases; the real sets carry the headline numbers. Its intervals are Wilson intervals that treat the frames as independent.

Blind-spot report. Every real benchmark set split by situation (by pose for SDT and by number of people elsewhere), and every stress case. Anything under 90% exact or 97% within one is flagged.

14 Owl: results

Figure 13 and Table 15 compare the shipped pair, owl-4.9 and its crowd model, with the rule on all six real benchmark sets. On five of them Owl counts 82–96% of frames exactly, where the rule as deployed manages 9–77%, and is within one person on 96–100%. On the crowded classroom set, the hardest set in this paper, it is exact on 39% and within one on 74%, against the deployed rule's 33% exact, a margin whose interval includes zero (−5.5 to 17.6 points).

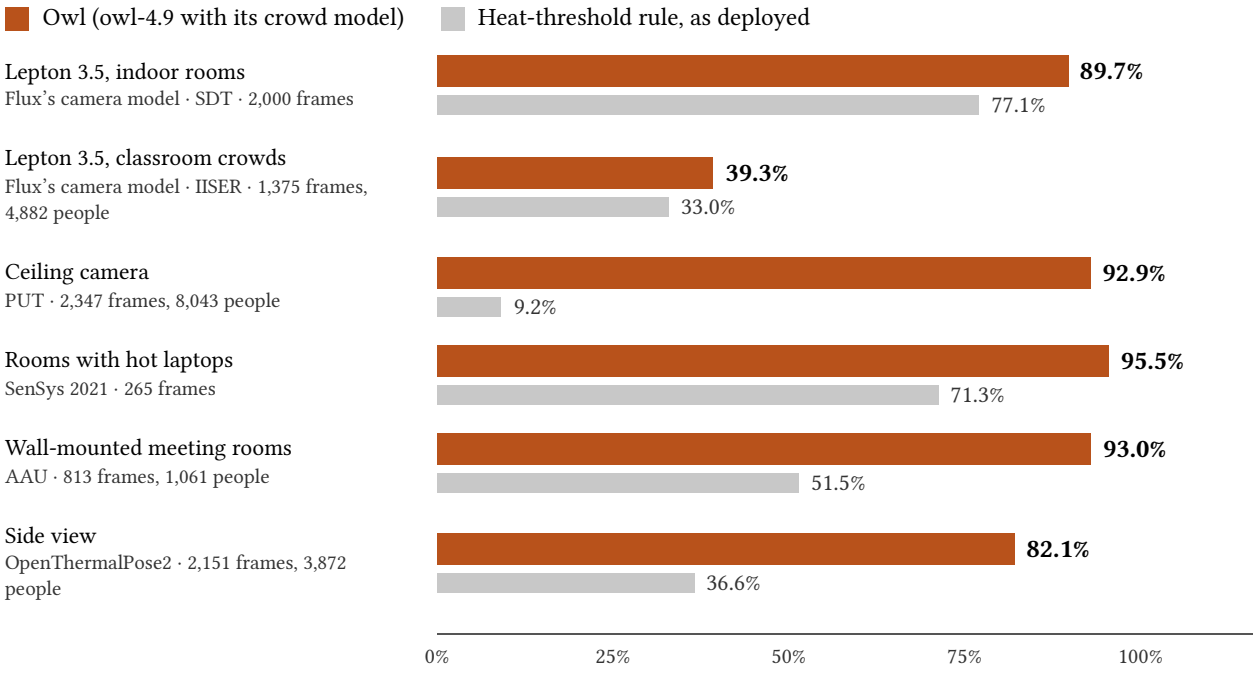


Figure 13. Frames counted exactly right on six real benchmark sets that no model trained on, owl-4.9 with its crowd model against the heat-threshold rule as deployed. The first two sets come from the FLIR Lepton 3.5, the camera model Flux uses. The rule is scored on identical frames; Table 15 adds intervals and the rule tuned to each set.

Table 15. owl-4.9 with its crowd model against the heat-threshold rule on the six real benchmark sets, scored on identical frames, with 95% block-bootstrap intervals for Owl's exact rate (a Wilson interval for the stress mean). Percent of frames. The last column is the rule with its bin, margin and minimum cluster size tuned separately on each set's validation split.

Benchmark set	Owl exact	Owl within one	Rule exact, as deployed	Rule exact, tuned per set
Lepton 3.5, indoor rooms (SDT)	89.7 (85.4–93.4)	100	77.1	93.7
Lepton 3.5, classroom crowds (IISER)	39.3 (31.2–47.7)	73.9	33.0	40.2
Ceiling camera (PUT)	92.9 (86.9–97.2)	99.7	9.2	29.4
Rooms with hot laptops (SenSys 2021)	95.5 (91.7–98.2)	100	71.3	82.6
Wall-mounted meeting rooms (AAU)	93.0 (86.4–98.0)	99.1	51.5	75.5
Side view (OpenThermalPose2)	82.1 (74.4–89.2)	96.3	36.6	60.6
41-case stress test, mean	94.7 (94.1–95.2)	99.3	–	–

A fairer rule. The deployed rule's settings were tuned on a bench, so I also tuned them on each set's validation split. Tuned this way, separately for each set, the rule matches or beats Owl on the two sets from Flux's camera model: 93.7% against 89.7% on SDT, whose frames hold one person or none and where most of Owl's errors are missed people (149 of 206), and 40.2% against 39.3% exact (84.6% against 73.9% within one) on IISER. On the other four sets it trails Owl by 13 to 64 points. The sensor's learned background mask made the tuned rule worse on every video set, because people who sit still fade into it. On the evidence here, a well-tuned threshold is a strong baseline on one camera in simple scenes; Owl's advantage is one model across cameras and views, and the warm objects that fool a threshold.

14.1 From owl-1 to owl-4.9

Owl went through 16 named versions in four days. Figure 14 follows the five original real benchmark sets across the versions that shipped. The jump between owl-2 and owl-3 is Findings 1 and 2 (Section 15): the same network with real frames mixed in. The dip on the meeting rooms at owl-4.7 and its recovery at owl-4.8 are a data-mix trade-off that Figure 17 isolates.

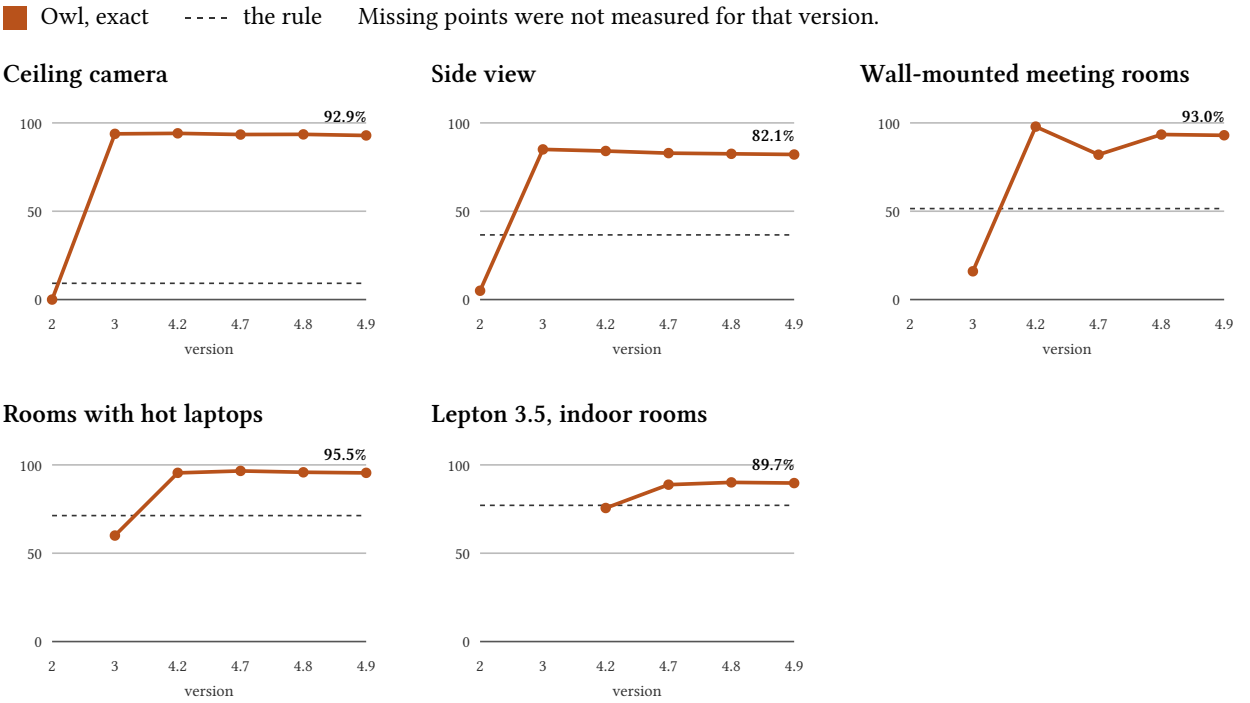


Figure 14. Exact counts on each real benchmark set's test split by version. owl-2 learned only from generated frames; owl-3 added real frames; owl-4.2 was the first shipped with meeting rooms, laptop rooms and measured physics; owl-4.7 and owl-4.8 added frames from the Lepton 3.5; owl-4.9 added real classroom crowds. owl-3's side-view figure includes the background fix of Finding 3 (74% before it).

14.2 Like-for-like comparison of the last two versions

When owl-4.9 was trained I re-graded owl-4.8 and owl-4.9 on the full test splits with identical code (Table 16). On the five earlier sets no exact-count rate moved by more than 0.6 points and no within-one rate by more than 0.8. On the classroom crowds the error fell by more than three quarters.

Table 16. owl-4.8 and owl-4.9, each with its crowd model, graded together on full test splits. Exact / within one, in percent. The owl-4.8 stress mean comes from a run before each case's seed was fixed; owl-4.9's is the seeded rerun (95% interval 94.1 to 95.2).

Test set	Rule exact	owl-4.8	owl-4.9
Lepton 3.5, indoor rooms	77.1	90.1 / 100	89.7 / 100
Lepton 3.5, classroom crowds	33.0	11.5 / 20.6	**39.3 / 73.9**
Ceiling camera	9.2	93.5 / 99.7	92.9 / 99.7
Rooms with hot laptops	71.3	95.8 / 100	95.5 / 100
Wall-mounted meeting rooms	51.5	93.4 / 99.9	93.0 / 99.1
Side view	36.6	82.5 / 96.6	82.1 / 96.3
41-case stress test, mean exact	–	94.5	94.7

14.3 Crowds on the Lepton 3.5

The IISER classroom set is the only real crowd I have from the Lepton 3.5: rows of 12 to 15 students, up to ten with a visible face, each row partly hidden behind the one in front. owl-4.8 had never seen anything like it and overcounted badly, averaging 8.1 people on frames that held 3.6. Figure 15 separates the two halves of the fix. Training the main model on the classroom frames lifted exact counts from 11.5% to 28.6%. Retraining the crowd model as well took them to 39.3%, and on frames with six to ten labelled people moved within-one from 3.4% to 96.1%.

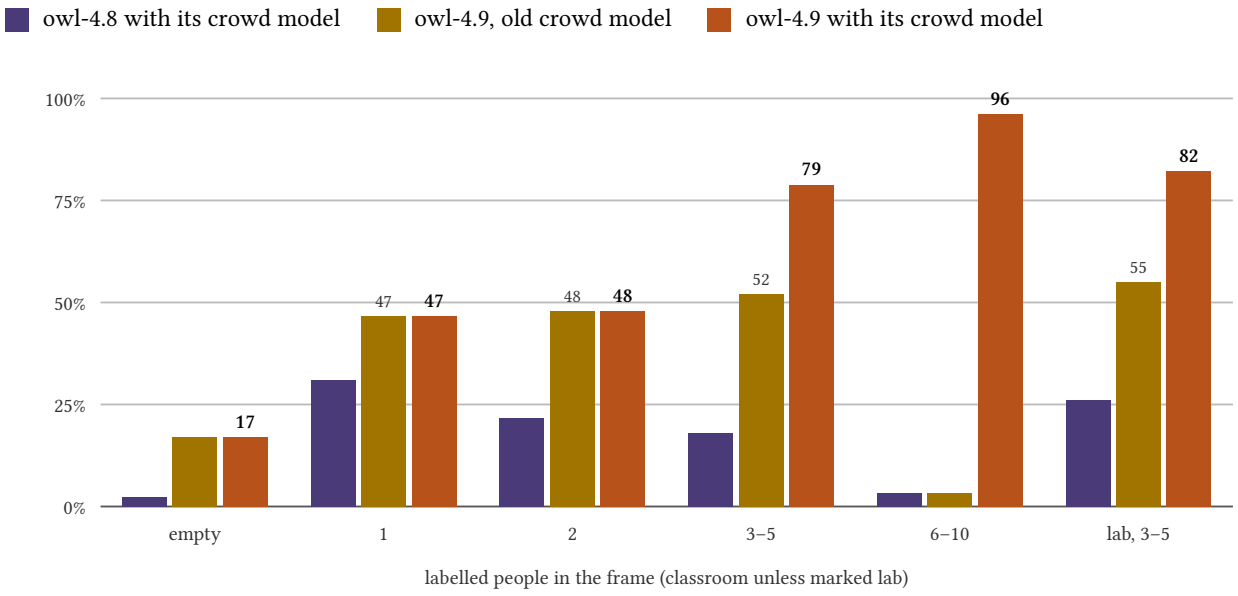


Figure 15. Frames within one person of the truth on IISER's held-out classroom and lab frames, by how many people were labelled in the frame. The crowd model's contribution is the difference between the gold and orange bars; it decides the six-to-ten bucket almost on its own.

By crowd size the classroom result splits in two. On frames with six to ten labelled people Owl is within one 96.1% of the time, against the deployed rule's 4.5%. On frames labelled with one or two people the rule is far better (37.9% and 72.7% exact, against Owl's 14.6% and 22.4%), and on the 88 frames labelled empty Owl reports 2.6 people on average. All of these sparse frames fall in IISER's test split, so Owl never trained on a sparse classroom. I inspected a sample: some hold people facing away or half in frame, whom the face labels omit, and in some Owl counts warm seat backs. A head-count relabelling of these frames will separate the two.

When the crowd model steps in is a single setting, chosen like the others on the classroom test frames and the stress test. Handing over earlier was worse on both (Table 17), so the sensor keeps five.

Table 17. When to ask the crowd model, owl-4.9 pair. The stress means and blind-spot counts come from runs before each case's seed was fixed, with about ±5 points of noise per case; the seeded rerun of the shipped setting gives a stress mean of 94.7%.

Crowd model asked from	Classroom exact	Within one	Stress mean	Blind spots
3 people	36.9	72.3	94.3	12
4 people	37.5	73.1	94.6	10
5 people (shipped)	**39.3**	**73.9**	**94.9**	11

14.4 Hard cases

Figure 16 shows all 41 stress cases for the shipped model, each generated from a seed fixed per case. Every warm object that fools the rule by design is counted correctly on 98 to 100% of stress frames, where the deployed rule manages 14% to 85%: a hand at the lens, a hot mug, a monitor, a radiator, a pet, a reflection, warm footprints, a candle. Four cases fall under 90% exact: people lying flat (87.3%), a bare torso filling the view at the lens (78.0%), one person behind another (60.7%, but 99.3% within one) and a crowd of ten, the weakest, at 31.3% exact (95% interval 24.5 to 39.1) and 75.3% within one. Stress figures for earlier versions in this paper come from runs before the seeds were fixed and carry about ±5 points of noise per case.

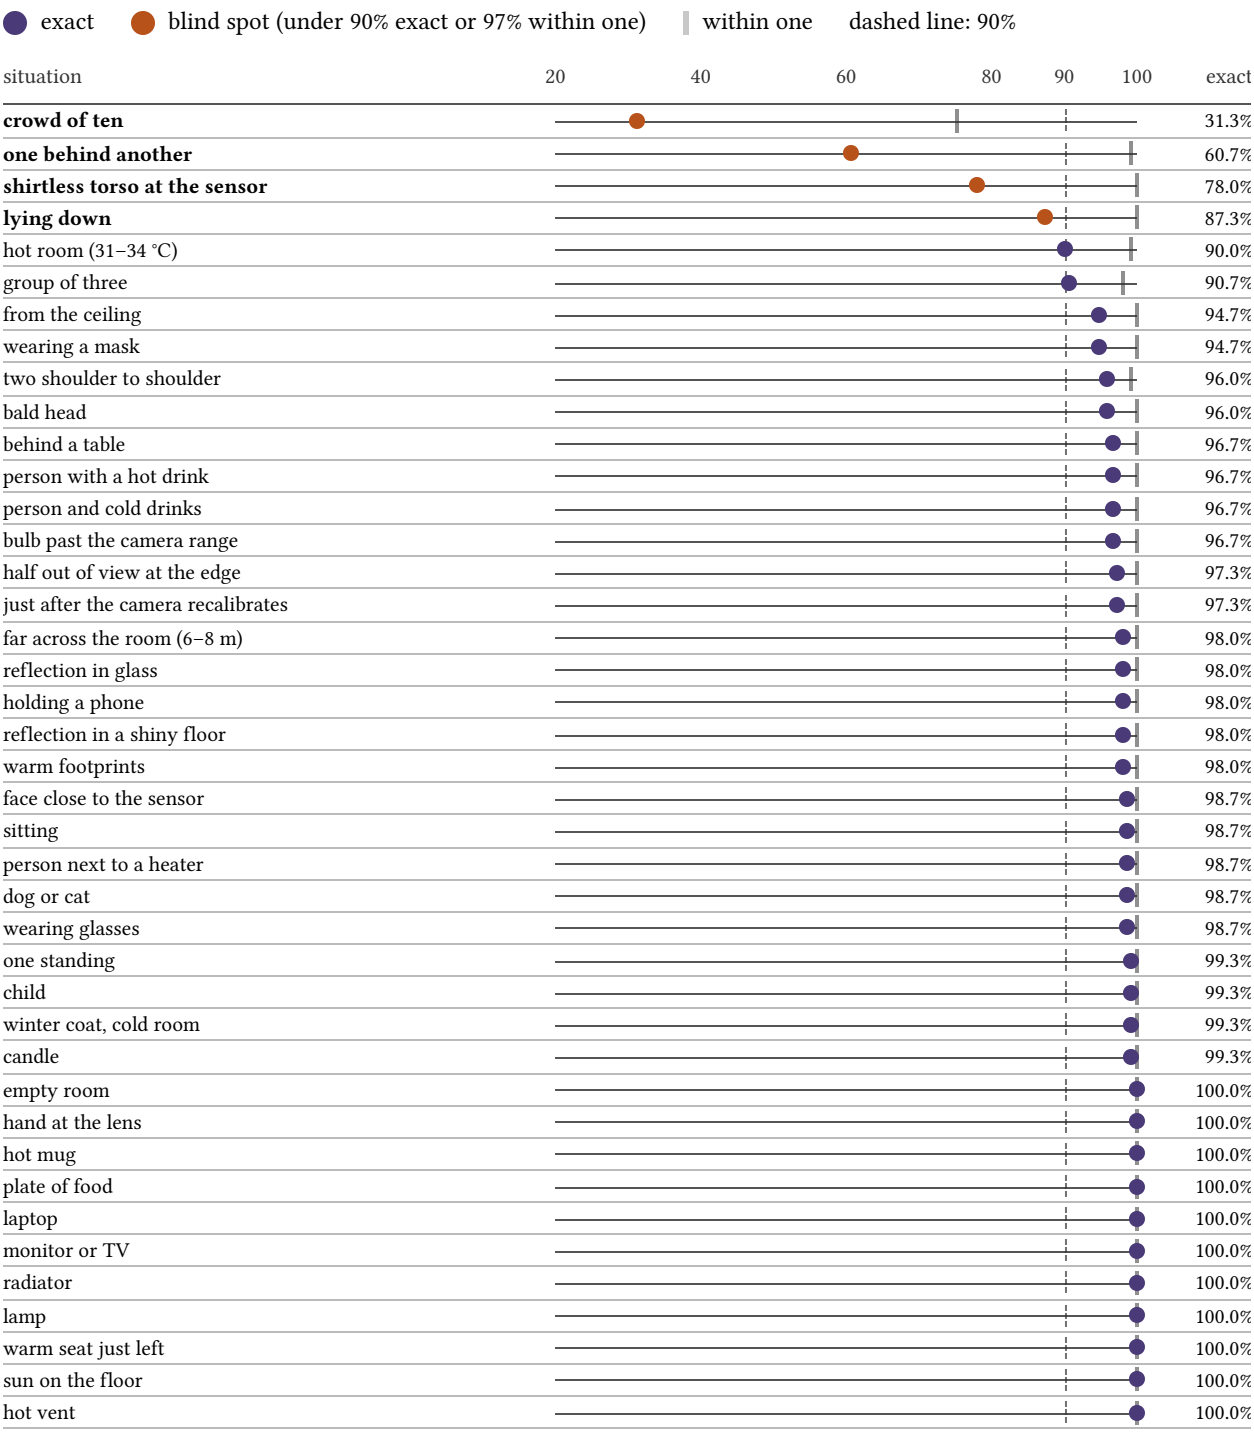


Figure 16. The 41-case stress test for owl-4.9 with its crowd model, 150 generated frames per case from a seed fixed per case, sorted from hardest to easiest. Mean 94.7% exact (95% interval 94.1 to 95.2) and 99.3% within one.

14.5 Naming everything else

On 3,000 held-out generated frames, owl-4.9 finds the nine other kinds as Table 18 shows. These names are shown on the sensor's screen to explain why something was not counted; they are never counted and never sent.

Table 18. Object naming on held-out generated frames. Found: share of real objects with a peak. Right when named: share of named peaks that were that kind.

Kind	Objects	Found	Right when named
person	10,461	74.9%	89.7%
hand	682	69.5%	84.3%
pet	617	67.1%	92.4%
hot drink	836	84.2%	83.0%
food	441	81.2%	91.6%
laptop	768	68.2%	79.8%
screen	491	60.5%	80.7%
heater	794	78.7%	85.5%
lamp	316	77.5%	79.0%
warm seat	586	58.7%	83.1%

The same 3,000 frames are deliberately harder than real rooms, with crowds of up to ten drawn often. On them owl-4.9 counts 59.3% exactly and 84.3% within one, against the rule's 20.2% and 49.9%; empty rooms are right 92.3% of the time and single people 80.1%. Because the generator grew harder with each version, these synthetic figures are not comparable across versions; each design finding names the evidence it rests on, whether real benchmark sets, generated stress cases, validation curves or the bench.

14.6 Speed and size

On the sensor, owl-4.8 and owl-4.9 take 14–16 ms per frame on one core of the Pi 5, measured by the sensor's own self test on its own camera. That is fast enough to read every frame the camera sends. An 8-bit quantised export runs 2.1 times faster at a quarter of the size, but it cost accuracy (Finding 11), so the sensor runs full precision.

15 Owl: findings

Every change to Owl was kept or thrown out on the real benchmark sets' test splits, which is why those sets are development benchmarks (Section 13). These are the results that shaped it, in the order I found them, including the alternatives that were tested and set aside. Stress-test figures in these findings come from runs before each case's seed was fixed and carry about ±5 points of noise per case.

F1 ***Generated frames alone fail on real footage***

owl-2 counted 80% of generated frames exactly and found 73–94% of the warm objects. On real footage it counted 0% of the ceiling frames and 5% of the side-view frames exactly. It had only seen clean drawn shapes and read the texture of a real room as people.

F2 ***A little real data fixes most of it***

A 500-step test with real frames mixed in went from 0% to 81% exact on the real ceiling set. The full run, owl-3, reached 94% on the ceiling set and 74% on the side view. From then on, generated frames also gained real-looking texture and low-resolution blocking.

F3 ***Judge warmth against the cool end of the frame***

The physical check compared each person with the room, taken as the frame's median. When a person fills the view, the median is the person, and the check threw out the person the model had just found. Taking the room as the 10th percentile lifted the real side view from 74% to 85% exact and 86% to 98% within one, hot rooms from 47% to 68%, and a bare torso at the sensor from 16% to 30%, with the ceiling set and every false-alarm case unchanged.

F4 ***Physics beats tricks***

Two augmentations that seemed sensible, skewing the absolute-temperature channel and drawing real group frames three times as often, changed nothing on the ceiling and side-view sets and cost a few points on unseen rooms. Replacing hand-picked temperatures with measured skin and clothing physics lifted the unseen meeting rooms from 24% to 33% exact and from 58% to 85% within one.

F5 ***A bigger network did not help***

Widening the network from 24 to 32 base channels tied the smaller one (94.2% and 73.9% against 93.8% and 74.2% on the ceiling and side view). Accuracy was limited by data, which set the direction for everything after.

F6 ***More of a hard case can cost more than it gains***

owl-4.4 drew more people lying down and more bare torsos and bodies at the sensor, the two single-person cases still under 90%. It lost the real meeting rooms, falling from 97.9% to 60.9% exact, with no measurable gain on the cases it targeted. The change was reverted.

F7 ***A sharp target at a guessed point teaches the model to miss***

A finer 80 × 60 output grid (owl-5) keeps heads a few pixels apart as two people. But on the datasets whose head points are estimated from a body box, a sharp target at the guess taught it to miss people: the meeting rooms fell from 97.9% to 67.2%. Spreading those targets over the pixels the head could occupy fixed the cause, and became standard for every later version.

F8 ***Density counts crowds that peaks merge***

Used as a second look in crowds, the finer-grid model lifted a crowd of ten from 41% to 63% within one. A density head, whose total does not fall when heads touch, lifted it to 82% and put one person behind another at 100% within one, with all four real benchmark sets unchanged. It became the crowd model (owl-4.3).

F9 ***Data shares trade off, and the right data ends the trade***

Adding real Lepton frames (SDT) improved results on Flux's camera model, but every way of making room for them hurt the meeting rooms: cutting their share to 15% cost 25 points, and owl-4.6 gained 13 points on the Lepton while losing 14 on the meeting rooms. owl-4.7 chose the camera Flux uses (88.8% against owl-4.2's 75.6%). Adding SDT's rendered Lepton-style frames in owl-4.8 improved both at once (Figure 17).

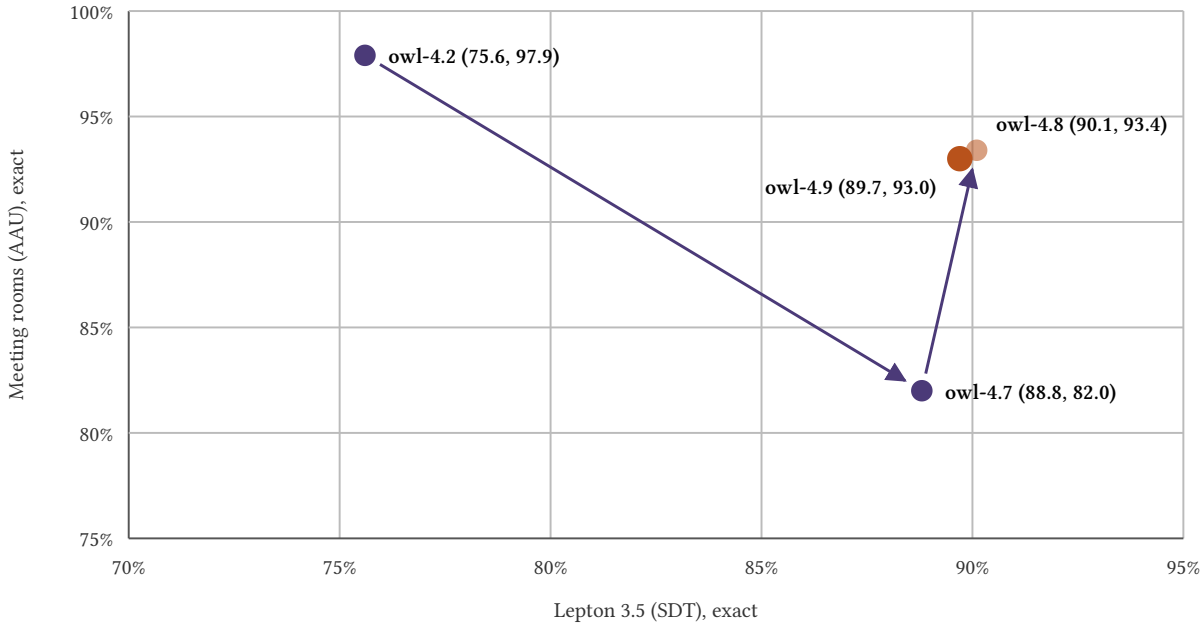


Figure 17. The trade between the Lepton 3.5 frames (SDT, horizontal) and the wall-mounted meeting rooms (AAU, vertical). owl-4.7 bought the camera with the meeting rooms; owl-4.8, trained with rendered Lepton frames as well, recovered the meeting rooms without giving the camera back. owl-4.9 differs from owl-4.8 by under one point on both sets.

F10 ***Validation scores do not compare across versions***

When training drew group frames three times as often, the validation draw reused that weighting by accident, and owl-3.2 looked worse than owl-3 on what was really a harder mix. Validation is still resampled with each run's own source weights, so a validation score chooses a checkpoint within its run and is never compared across versions.

F11 ***Quantising cost more than it saved***

An 8-bit export ran 2.1 times faster at a quarter of the size, but cut owl-3's exact count on the real ceiling set from 94% to 86%, and to 84% when calibrated on real frames. At 14–16 ms per frame in full precision, speed was never the constraint, so the sensor keeps full precision.

F12 ***A newer base can make a worse crowd model***

Retraining the crowd model starting from owl-4.8 instead of its own earlier weights dropped a crowd of ten to 18.7% exact. The older crowd model stayed until real crowd data arrived.

F13 ***A real room shows what no test set does***

On the sensor in a real room, owl-3 named a pet in a room with none, sometimes counted a monitor, and missed a shirtless torso close to the camera. Three changes answered them: higher naming bars for rarer kinds, dropping a person where a screen-like object is more confident, and bare skin and bodies filling the view in the generated frames. On the stress test a shirtless torso at the sensor went from 16% under the first owl-3 decoder (30% after Finding 3) to 78% under owl-4.9 with the seeds fixed.

F14 ***Real crowds on the Lepton 3.5 were the missing data***

After owl-4.8, training on the same data had stopped helping. Only after IISER Bhopal licensed its Lepton 3.5 classroom crowds to me did the crowd numbers move: 11.5% to 39.3% exact, an average error of 4.66 people to 1.03, and rooms of six to ten from 3.4% to 96.1% within one, while on the other benchmark sets no exact-count rate moved by more than 0.6 points and no within-one rate by more than 0.8.

F15 ***Validation stopped improving early in the run***

owl-4.9's real validation score stopped rising about a quarter of the way through its run (Figure 18). The rest of the run brought no sustained gain, so the next experiment is new data: more real crowds and side views from this camera.

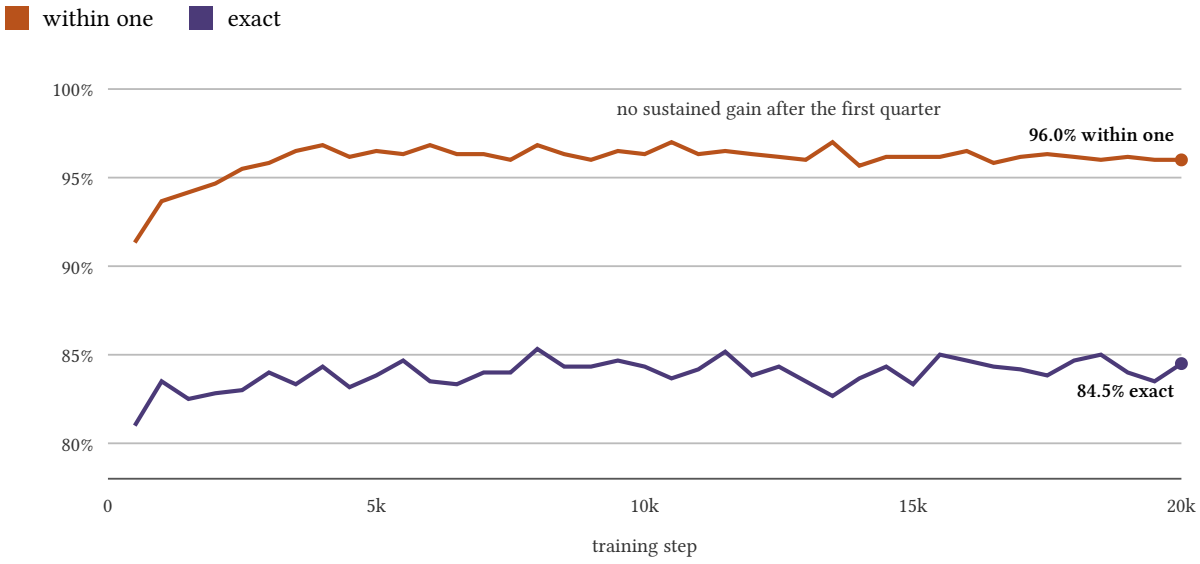


Figure 18. owl-4.9's real validation score every 500 steps over its 20,000-step run. Both curves are flat after about 5,000 steps. The validation mix includes the new classroom crowds, so these values are lower than earlier versions' and are not comparable with them.

16 Privacy by design

Flux is built to tell people how busy a place is. It must never become a camera that watches them. Owl reads each thermal frame on the device and keeps no history of frames.

- **Kinds, never who.** Owl's outputs are ten kinds of thing. It has no identity output, no face class and no face features. IISER labels faces, and I use each face box only as the point where a head is; Owl is trained to locate people, never to predict who they are.
- **No memory between frames.** Owl reads each thermal frame on its own and keeps no track of anyone. The doorway counter follows a crossing only for the moment it takes to tell an arrival from a departure.
- **Numbers are all that leave.** The network's points and boxes exist only in memory. Boxes are drawn only on units with a screen and are never sent. What leaves the sensor is a handful of numbers: people counts and a sound level, never an image.
- **Enforced by tests.** Tests in the sensor's suite fail the build if Owl gains an identity output or carries anything between frames, or if an image, a point or a box leaves the device.
- **Data used as licensed.** Every dataset is used as its licence or permission allows. None of their images are redistributed or published, and none are shipped with the model. IISER's licence also forbids using its data to identify anyone, and Owl has no identity output.

The model files are kept private, like Flock's crowd-forecasting model. The generator, network and training and grading code are public. Without a model on disk the sensor counts with the rule, and says so.

17 Future work on the counter

The clearest next step is data. A request to use a further thermal crowd dataset is outstanding, and the gap that matters most is real radiometric crowds from the Lepton 3.5 in venues. A study in a working venue, with hand counts as ground truth, would replace the stress test's estimate of real-world accuracy with a measurement. Before that, a sealed test of frames that no version has seen will be scored once, as set out under What comes next. Flux is built to carry a doorway counter, an 8 × 8 time-of-flight sensor that tells arrivals from departures; it first ran on the bench in October 2026 and has not yet counted a real doorway. Combining its running total with Owl's count of people in view, and estimating dwell time with Little's law [28], would answer how many people are inside as well as how many are visible, without identifying anyone. Finally, each crowd decision currently rests on one frame; agreeing a count across consecutive frames, without following individuals, may steady crowded readings further.

18 Conclusions on the counter

Owl shows that a model small enough for one core of a Raspberry Pi can count people on a 160 × 120 thermal camera across cameras, views and warm distractors, without recording new footage of anyone to train it and without the sensor keeping a frame after it is counted. Generated frames gave it a vocabulary of warm things that the rule confused with people; licensed real frames taught it what real

rooms look like; and each real dataset, once added, closed the gap it was chosen for. On five of the six real benchmark sets, owl-4.9 counts 82 to 96% of frames exactly, where the deployed rule manages 9 to 77%. A rule tuned to each set matches it on the two sets from Flux's own camera, and on real classroom crowds Owl misses by one person on average, where its predecessor missed by almost five. Because the benchmark sets also guided its development, the next measurements are a sealed test scored once, more real footage of crowds from the camera Flux uses, and a first measurement in a working venue.

19 How it was built

Both studies were built in seven months, from the first collector in March 2026 to the sensor's doorway counter in October. This section records the process: what was built, in what order, each model version and whether it passed its gate, the alternatives tested and the decisions they produced, the parity tests that hold the system together, and the hardware. Table 19 gives the timeline. Dates are Eastern time.

19.1 The forecast: corpus and first models (March to May 2026)

The first collector was committed on 10 March 2026 and widened to 34 cities the same day. From the start it gathered two kinds of row from the busyness provider: weekly rows, one typical week per venue of 168 hourly values on the provider's 0 to 100 scale, and realtime rows, one reading per venue-hour on a real date. Event listings were attached the same week, and a model pipeline with an ONNX inference service followed on 14 March. In early May the collector reached the 13 of its 34 cities it had not yet tried: a probe of 100 venues in each found 996 of 1,300 (76.6%), and the venues found entered the corpus. Beijing, where 37 of 100 matched, stayed thin and holds no rows in v2.6.0's training frame, so 30 cities train the model and three are held out (Section 4.5). The last spring reading is dated 18 May 2026, and by then the corpus held 3,912,357 rows across 34,784 venues in 34 cities (Section 4.2). That spring corpus trained the shipped delta model; the September rehearsal trained on live readings, and Owl on separate thermal datasets.

The first models predicted busyness directly, with the venue's popular-times value among their inputs. An 80-feature model of this kind scored far higher on validation than on held-out cities: it had learned to copy the curve, and coordinates acted as a lookup table for the five cities it was trained on. That result set three design decisions that still stand. Coordinates were dropped as features. Cross-validation holds out a whole city at a time. And the label became the *delta*, how far the hour departs from the venue's own curve, with the served score the curve plus a clamped correction. On 10 May, v2.2.1 read the curve from the same table serving reads, the project's first train/serve parity fix. No accuracy figure from before August is used as evidence, because the corpus was later repaired (Section 4.2).

19.2 An end-to-end audit, and v2.6.0 (August 2026)

In August I audited the whole pipeline end to end. The audit found that the model had never loaded in production, because the runtime package that reads an ONNX file was missing from the backend's dependencies, so the rule engine had answered every request; and that the collector had written no rows since 19 May while reporting success, because its provenance check compared zero rows written with zero rows labelled. Both checks were rewritten to fail loudly. On 12 August a training ladder ran:

v2.3 trained on the rows production serves, realtime rows with a usable curve, because weekly rows have a delta of zero by construction and had taught earlier models to shrink every correction. v2.3.0 to v2.4 failed their gates, and v2.5.0-starling shipped that afternoon. The next six days were spent on the data and the pipeline:

- *Train/serve parity* (14 August). Serving was not building the vector training built: live predictions had been computed at coordinates 0, 0, inference read a venue's whole type list where training read the first three, and the event-size estimator had drifted from the one that labelled the corpus. Model start-up now fails closed on a parity break.
- *The clock* (15 August). The weekly collector had stored the provider's array index as the hour; the provider's day runs from 06:00, so slot 18 was the venue's midnight. A migration corrected 3,536,360 rows, and a per-row clock column now refuses an undeclared row.
- *The audit* (15 August). A read-only pre-retrain audit listed 21 findings, 8 of them blocking, among them 491,100 duplicate rows, an impossible month on 62.9% of rows and a gate that scored rows production never serves. All eight were fixed before the next run.
- *Provenance* (15 August). The 457,402 spring realtime rows cannot be split into live readings and provider forecasts; two attempts to recover the split failed against null controls. Those rows stay marked unknown, and since 4 September they are excluded from training by default.
- *A category leak* (16 August). Two category features had been fitted on the whole frame, so a held-out city helped build its own scores. They are now refitted inside each fold.

The first retrain in which every stage had been checked shipped as v2.6.0-starling on 18 August, with the result of Table 6: MAE 29.42 against 31.48 for the curve and 30.77 for v2.5.0 on 67,249 rows in three cities it never trained on. The gate's old within-10 floor of 29.2% had been derived before the clock fix, so it was replaced by a measurement, the incumbent's 19.3% on the same rows, which v2.6.0 beat at 20.7%. Two copies were trained, and the CPU artifact shipped over the GPU one, which scored within 0.02 MAE, because the CPU run reproduces bit for bit.

19.3 Alternatives tested, and the decisions they produced (August to September 2026)

After v2.6.0, every proposed change was measured against a gate before it could ship, and each result is recorded so the experiment is not repeated.

- *Calendar features* (19 August). Sweeping the month across the year moved v2.6.0's raw correction by 27.31 points on average, against 5.92 for sweeping the hour: a one-spring corpus stores collection drift in the calendar. v2.7, without month and season, failed its gate (MAE 29.89, within-10 19.9%), so v2.6.0 stayed, and the October retrain drops the month family with live readings to replace it (Section 7.2).
- *Calibration* (20 August). About 40 post-hoc corrections were fitted on the earliest 30% of dates and scored forward on 46,101 rows. On that forward split a quantile display map gained 8.40 points of within-10 for 3.15 MAE (Section 6.5) and was armed on 28 August; September's live readings later showed it cost within-one accuracy, and it left the served number on 26 September (Section 6.5).

- *Hour ranking* (20 August). Within one venue-night the model ordered two hours correctly 62.7% of the time and the curve 63.1%, over 31,498 pairs, so the app's "go later" advice reads the curve and stays silent below a 10-point gap.
- *Event provenance* (20 August). 22 of 34 cities held no row with a nearby event because their events were never collected. A new provenance column now separates a measured absence from a missing lookup.
- *Game nights* (30 to 31 August). Games of seven teams became six features. Game nights are 7.1 points busier in the labels, but the spring fit window held about 19 days of games, and the features did not move within-10. They wait behind a switch for the October corpus.
- *The two-head model* (5 September). Held to the proven-live labels then available, it lost to the shipped model, and a model given 304,000 more rows of the same kind lost to one trained without them, in a single run without an interval (Section 6.5). Both results pointed to the venue's own recent readings as the input to build.
- *Forty-two levers* (6 September). Six research passes proposed 42 changes, each attacked by a reviewer told to refute it; seven survived. The one shipped was the per-venue trailing offset (Section 6.5).
- *The curve's range* (6 to 8 September). The poor calibration of Section 3.1 invited a slope correction; it bought MAE and spent within-10, so the system keeps the per-venue intercept (Section 6.5). The same week measured what coverage is worth (Section 6.4), which made adding venues with a curve the collector's first job.

19.4 Live collection (September 2026)

Live collection restarted on 1 September for about 1,400 venues in Philadelphia and the Lehigh Valley, and every row now carries a provenance label. Pacing settled at one call a second after the provider's rate guard throttled the first day's 600 calls a minute. The sweep went hourly from 20:08 on 3 September (00:08 UTC on the 4th) and now runs at seven minutes past each hour, and an open-hours filter stopped calling venues that their own curve says are shut, cutting the mean sweep from 1,414 calls to 993. From 5 September, Miami, one of the three held-out cities, is read at two fixed hours a day, so the gate has live readings from a city training never sees. From 6 September the venues users actually open lead every sweep, and from 25 September calls run in parallel, up to eight in flight, so the full sweep finishes inside the hour. The proven-live pool grew from 4,223 readings on 4 September to 9,339 over 1,228 venues on 6 September.

19.5 The switch to curve, offset and nowcast (25 and 26 September 2026)

On 25 September a replay of the serving code over every live reading in the 8 September export measured, for the first time, what the card is scored on: whether the served number lands in the same crowd level as the live reading. As served, v2.6.0 with the display map landed within one level 57.4% of the time on 7,920 readings, and a reading one hour old predicted the next within one level 92.0% of the time. Eleven causes of miss were ranked by the points each could recover, and the two largest needed no retrain. Two serving switches followed: the curve plus 0.85 of the trailing offset, and the nowcast of

Equation (2), with weights fitted on 1 to 5 September and scored on 6 to 8 September. The switches went through six adversarial review rounds, one of which caught that the offset could see the hour it was scoring; the fix restricted it to readings strictly before the hour. They went live in production at 18:55 on 26 September, on the three-day result in the last row of Table 4. On 25 and 26 September the October retrain's rules were committed (Section 7.2), and on 26 September the retrain was rehearsed (Section 6.6). On 8 October a fresh read-only export scored the switches on every live reading since 9 September: 54.9% exact on 119,939 readings over 30 dates, against 56.0% on the three days they were shipped on (Section 6.1). On 13 October I scored the month against persistence baselines: a rule that repeats an hour-old reading, or else serves curve plus offset, matched the served number, because the two-hour weight fitted in early September no longer helped (Sections 6.1 and 6.3). Production was left as it was, so the October gate compares against what was registered.

19.6 The parity tests

The August audit found faults in collection, provenance, model loading and train/serve parity. The parity faults were the hardest to see, so the safeguards built since are tests that training and serving agree:

- the ONNX export verifies feature order, and serving refuses an artifact that lists a feature it cannot build;
- a holdout with zero cities must reproduce the shipped category features bit for bit before any fold is fitted, with a negative control that rebuilds them the old way and must be caught;
- the curve smoothing, the neighbour features, the live features, the game features and 62 calendar, weather and event columns are each computed by both paths over shared fixtures, with zero disagreements allowed;
- the collector writes the event values serving computes, checked over 400 listings;
- the band replay must publish what the real prediction function publishes on every row;
- a leak test changes every reading at or after a row's hour, including its own label, and requires none of the row's inputs to move;
- the serving responses of v2.6.0 are pinned under every switch combination.

19.7 Flux, the sensor (May to October 2026)

The sensor began as backend tables and an endpoint on 2 May 2026, tested end to end before any device existed. The demo unit's design was fixed on 4 May, built from about $1,000 of parts: a Raspberry Pi 5, a FLIR Lepton 3.5 on a PureThermal 3 board, a 4G board, an analog microphone, an infrared beam and a battery. The Lepton was chosen over the 32 × 24 MLX90640 for its 160 × 120 radiometric image; moving the sensor program onto it on 26 August took the thermal path from 768 pixels to 19,200. The first bench day on a Pi 5 was 6 September: the camera came up on the first try, and pooling the image into 4 × 4 cells stopped a bare head and a clothed torso from reading as separate people. On 21 September the enclosure was designed in OpenSCAD and the device split in two, a base with the Pi, battery and screen, and a wall head with the sensors looking down at a doorway. On 7 October the analog

microphone gave way to an I2C decibel meter that measures A-weighted sound level on its own processor, so the Pi reads a number and never audio. On 8 October the base and the wall head were merged into one box, because the doorway sensor and the decibel meter need short I2C wires.

19.8 Owl, the training chain (26 to 29 September 2026)

Owl replaced the heat-threshold rule in four days, through 16 named versions, 19 full training runs and five short trials (581,700 steps, 37.2 million frames) on one RTX 5080. owl-1 (26 September) learned from generated thermal scenes alone; owl-2 and owl-3 (27 September) showed that real frames must be mixed in; owl-4.2 to owl-4.8 searched for the right mix of data; owl-4.9 (29 September) added real classroom crowds. Findings 1 to 15 (Section 15) record the design decisions this chain produced and the evidence behind each: the real benchmark sets' test splits, generated stress cases, validation curves and the bench. Because versions and settings were kept or rejected on those test splits, the six real sets are development benchmarks (Section 13).

Every real frame is used under a commercial licence or with written permission. Requests went out on 27 September; TU Wien granted SDT on 28 September, and IISER Bhopal licensed its Lepton 3.5 classroom crowds on 29 September. A dataset whose authors had not replied was not used, and another was left out because its labels carry a copyleft licence. The shipped pair, owl-4.9 and its crowd model, draws 60% of its training frames from licensed datasets (SDT's rendered Lepton frames among them) and 40% from the generator.

19.9 The doorway counter (September and October 2026)

The crossing sensor went through three designs. The infrared break-beam of the May plan became a one-sided infrared proximity module on 21 September, with one part to mount and nothing to align, and on 25 September a VL53L8CX time-of-flight sensor that looks down at the doorway as an 8 × 8 grid of distances and tells arrivals from departures. Its driver was written on 26 September, before the part arrived, as a port of the manufacturer's driver onto the Python standard library, together with an on-device occupancy estimate that combines the door's running total with the camera's count as a floor and estimates dwell time with Little's law [28]. On 7 October the counter ran on the board for the first time, reading 265 frames.

Table 19. Timeline of both studies, 2026. F: the forecast; C: the counter.

Date	Track	Milestone	Result
10 Mar	F	First collector; 34 cities	Weekly curves and readings begin
14 Mar	F	Model pipeline and ONNX inference service	
Mar to Apr	F	Direct models with the curve as an input	Copied the curve; delta label adopted
2 May	C	Sensor tables and endpoint	Tested before any device existed
4 May	C	Demo unit designed; Lepton 3.5 chosen	About $1,000 of parts
10 May	F	v2.2.1 reads serving's curve table	First parity fix
18 May	F	Last spring reading	Corpus of 3,912,357 rows

Date	Track	Milestone	Result
12 Aug	F	Audit; v2.3 to v2.4 fail; v2.5.0 ships	Model loading and collector checks rebuilt
14 to 16 Aug	F	Parity sweep, clock fix, 21-finding audit, leak closed	3,536,360 rows corrected
18 Aug	F	v2.6.0-starling ships	MAE 29.42 against 31.48, 67,249 rows
19 Aug	F	v2.7 without month and season	Fails gate, MAE 29.89
20 Aug	F	About 40 calibrations; quantile map	Map +8.40 within-10, +3.15 MAE (46,101 forward rows)
26 Aug	C	Sensor program moves to the Lepton 3.5	768 to 19,200 pixels
30 Aug	F	Game-night features	Within-10 unchanged; held for October
1 Sep	F	Live collection restarts with provenance	One call a second
3 Sep	F	Hourly sweeps begin	
5 Sep	F	Two-head model	Two-head MAE +1.59
5 Sep	F	Miami holdout hours begin	Two runs a day
6 Sep	F	Demand-first sweeps; trailing offset ships	MAE 25.90 to 24.53
6 Sep	C	First bench day on a Pi 5	4 × 4 pooling
6 to 8 Sep	F	Curve range measured; slope refused	Slope 0.466; coverage worth 1.82 MAE
21 Sep	C	Enclosure; base and wall head	
25 Sep	F	Band replay; parallel sweeps; retrain rules committed	As served 57.4% within one
25 Sep	C	Time-of-flight doorway counter chosen	
26 Sep	F	Curve + offset + nowcast in production, 18:55	Exact level 32.2% to 56.0% on the three-day test
26 Sep	F	Retrain rehearsed with live features	79.7% within one; refused by the gate
26 Sep	C	owl-1, generated frames only	
27 Sep	C	owl-2 fails on real frames; owl-3 and owl-4.2	Real ceiling 0% to 94%
28 Sep	C	Crowd model; owl-4.7 and owl-4.8; SDT permission	Flux's camera model 75.6% to 90.1%
29 Sep	C	owl-4.9 and its crowd model ship	Crowd error 4.66 to 1.03 people
3 Oct	F	Provider's live coverage drops	Live readings fall to about 61% of the week before
6 Oct	F	Retrain set for 19 to 20 October	
7 Oct	C	Doorway counter's first frames; decibel meter	265 frames
8 Oct	F	Month-long test of the served number	54.9% exact on 119,939 readings, 30 dates
12 Oct	F	Provider returns no live readings from about 16:00	
13 Oct	F	Persistence baselines scored on the month	Hour-old reading, else curve + offset, matches production
13 Oct	C	Owl re-scored with intervals, a tuned rule and a seeded stress test; sealed test set aside	Stress mean 94.7%

20 Limitations

- **Relative readings, not headcounts.** Every forecast label is the busyness provider's relative reading on a per-venue 0 to 100 scale, an estimate made by the provider's own system. No hand count or sensor count has checked it, and a score of 70 means different crowds at different venues. What the forecast predicts is the provider's next live value. Consecutive-hour readings were identical 80.9% of the time in early September and about 45% to 48% from late September (Section 3.2), which may reflect how often the provider refreshes its value; part of the persistence result rests on reading an unchanged value twice.
- **The test window.** The served accuracy rests on 119,939 readings over 30 dates (9 September to 8 October) in Philadelphia, the Lehigh Valley and Miami: three of the 34 cities, in one season, and the nowcast's gain comes from the two Pennsylvania regions alone. The online correction went into production on a three-day test, before the five-date band gate a model artifact faces; the 30-day test confirms it, and the October test is the first pre-registered one. Persistence baselines match the served number, and the best of them was found on the same window (Section 6.1). On *Packed* readings the served number finds fewer than the configuration it replaced did (recall 45.3% against 49.2%, on 20,445 readings) but is right far more often when it says *Packed* (67.4% against 28.0%).
- **The delta model's training data and its holdout.** The model results of Section 6.2 rest on spring rows of unknown provenance, and its three held-out cities also scored the versions that failed before it, so they act as a validation set. On verified live readings the spring-trained model loses to the curve: over the 30-day test, MAE 30.94 as served (display map on) against 26.40 for the weekly curve, and in the 8 September export the model alone scored 27.52 against 26.42 for the curve on 7,920 readings. Ten of its 106 features are never split on (Section 5.1), latitude re-enters through daylight hours, and the learner comparison of Table 11 is untuned and confined to three cities.
- **Owl's benchmarks guided its development, and it has not been measured in a venue.** The six real sets were used to choose between versions and settings, so Figure 13 and Tables 15 and 16 are development results, and a sealed test has not yet been scored (Section 21.2). All six come from other people's rooms; IISER's truth counts visible faces; AAU's training and test splits share rooms and people. A rule tuned to each set matches Owl on the two Lepton sets. The stress test comes from the same generator as 40% of the training frames and was used to steer training. People hidden by people remain the hardest case (a side view of three to five people is exact on 40% of frames, 85.5% within one), lying and sitting people on the Lepton are exact on 81% and 83%, and on sparse classroom frames Owl overcounts. SDT's Lepton frames hold one person or none, and a study with hand counts in a working venue has not been done.
- **No Flux is installed at a venue.** The one unit is a demonstration unit, Flock's crowd number reads no sensor data, and the doorway counter has run on the bench alone.

21 What comes next: the retrain, and joining the forecast to the counter

Both studies serve one number in Flock: how busy a place will be when a group arrives. The forecast now delivers that number from the venue's own recent readings, and Owl counts people on a device small enough to mount in a doorway. This section sets out the next measurements; none of them has been run, and this paper reports no results from them.

21.1 The October retrain

Section 7.2 sets out the retrain, its pre-registered rules and its held-out window; I will update this paper with its results, whichever way the gate decides.

21.2 A sealed test for Owl

Owl's six real benchmark sets also guided its development (Section 13), so its next accuracy figure comes from frames kept apart. Two sets of frames have been put aside and will be scored once, after they are hand-counted. The first is 300 frames from the office test split of OpenThermalAction, a public MIT-licensed thermal set that no Owl version has seen, spread over 269 clips of one and two people. The second is the IISER test frames labelled with two or fewer faces, plus 100 others, recounted by head instead of by face. Owl's and the rule's counts on the first set were computed and stored before any hand count, so neither the model nor its settings can be adjusted to them.

21.3 Counts as ground truth

The forecast predicts busyness normalised per venue, so that 100 is the busiest that venue gets (Section 4.1). A score of 70 could mean a dozen people in a small bar or two hundred in a large one. A Flux unit in a venue that Flock already reads would give that venue an absolute scale: its headcount, paired with the busyness provider's live score for the same hours, would show how many people a given score means at that venue, and whether the five crowd levels match what a group sees at the door. Its doorway counter, once proven on a real door, would add arrivals and departures, so the sensor could estimate how many people are inside as well as how many are in view (Section 17).

It would also give the nowcast a fresher reading. The nowcast weights the latest reading by its age (Section 5.2), and over the test month a reading's value was spent within about an hour (Sections 3.2 and 6.3). A sensor that reports every 30 seconds would make that reading minutes old, where today's readings are an hour old at best and often several hours old or missing (Figure 8).

21.4 What would be measured first

- One unit, installed with the owner's agreement, at a venue the hourly collector already reads.
- Hand counts of the people in Owl's view, timed against its counts on several busy nights, to score Owl in that room before its counts are used as ground truth, and hand counts of arrivals and departures to score the doorway counter.
- Hour by hour over several weeks: the busyness provider's live score against Flux's count, the score per person, and how often the crowd level names what the count says.

- The October retrain, scored exactly as pre-registered. The counts are a separate test and change none of its rules.

Together these turn the two studies into one measurement of how busy a place is.

Acknowledgements

Owl would not count real rooms without the people who shared their data. The Computer Vision Lab at TU Wien (Christopher Pramerdorfer, Julian Strohmayer and Martin Kampel) granted permission to use the SDT dataset, real and rendered. Arijit Samal and Haroon R. Lone of IISER Bhopal licensed the Thermal Vision dataset and granted permission to use it. PUT Thermo Presence (PUT Vision Lab) and OpenThermalPose2 (ISSAI) are used under the MIT licence; the AAU VAP Trimodal dataset (Aalborg University) and the SenSys 2021 low-resolution set are used under Creative Commons Attribution 4.0. None of their images are included in Flux or in this paper.

Data sources. Venue busyness readings come from a commercial venue-busyness provider, which has no part in and does not endorse this work. Weather data is provided by OpenWeather (https://openweathermap.org/) under the Open Database License (ODbL). Sports schedules are from TheSportsDB. Event listings are from Ticketmaster. No dataset from these sources is released with this paper.

References

[1] J. Bollenbach, S. Neubig, A. Hein, R. Keller, and H. Krcmar. 2022. Using machine learning to predict POI occupancy to reduce overcrowding. In *INFORMATIK 2022* (LNI P-326). Gesellschaft für Informatik, Bonn, 393–408. https://doi.org/10.18420/inf2022_34

[2] J. Bollenbach, S. Neubig, A. Hein, R. Keller, and H. Krcmar. 2024. Enabling active visitor management: local, short-term occupancy prediction at a touristic point of interest. *Information Technology & Tourism* 26, 3 (2024), 521–552. https://doi.org/10.1007/s40558-024-00291-2

[3] S. Chang, E. Pierson, P. W. Koh, J. Gerardin, B. Redbird, D. Grusky, and J. Leskovec. 2021. Mobility network models of COVID-19 explain inequities and inform reopening. *Nature* 589, 7840 (2021), 82–87. https://doi.org/10.1038/s41586-020-2923-3

[4] T. Chen and C. Guestrin. 2016. XGBoost: a scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '16)*. ACM, 785–794. https://doi.org/10.1145/2939672.2939785

[5] J. B. Copas. 1983. Regression, prediction and shrinkage. *Journal of the Royal Statistical Society: Series B (Methodological)* 45, 3 (1983), 311–335. https://doi.org/10.1111/j.2517-6161.1983.tb01258.x

[6] F. Galton. 1886. Regression towards mediocrity in hereditary stature. *The Journal of the Anthropological Institute of Great Britain and Ireland* 15 (1886), 246–263. https://doi.org/10.2307/2841583

[7] A. Hajisafi, H. Lin, S. Shaham, H. Hu, M. D. Siampou, Y.-Y. Chiang, and C. Shahabi. 2023. Learning dynamic graphs from all contextual information for accurate point-of-interest visit forecasting. In *Proceedings of the 31st ACM International Conference on Advances in Geographic Information Systems (SIGSPATIAL '23)*. ACM, 1–12. https://doi.org/10.1145/3589132.3625567

[8] R. J. Hyndman and A. B. Koehler. 2006. Another look at measures of forecast accuracy. *International Journal of Forecasting* 22, 4 (2006), 679–688. https://doi.org/10.1016/j.ijforecast.2006.03.001

[9] S. Islam, D. Gandhi, J. Elarde, T. Anderson, A. Roess, T. F. Leslie, H. Kavak, and A. Züfle. 2021. Spatiotemporal prediction of foot traffic. In *Proceedings of the 5th ACM SIGSPATIAL International Workshop on Location-based Recommendations, Geosocial Networks and Geoadvertising (LocalRec '21)*. ACM, 1–8. https://doi.org/10.1145/3486183.3490997

[10] S. Kapoor and A. Narayanan. 2023. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns* 4, 9 (2023), 100804. https://doi.org/10.1016/j.patter.2023.100804

[11] M. Möhring, B. Keller, R. Schmidt, and S. Dacko. 2021. Google Popular Times: towards a better understanding of tourist customer patronage behavior. *Tourism Review* 76, 3 (2021), 533–569. https://doi.org/10.1108/TR-10-2018-0152

[12] D. S. Moore, W. I. Notz, and M. A. Fligner. 2013. *The Basic Practice of Statistics* (6th ed.). W. H. Freeman and Company, New York.

[13] ONNX Runtime developers. ONNX Runtime. Software, accessed October 2026. https://onnxruntime.ai/. Server: onnxruntime-node 1.29.0; sensor: the Debian python3-onnxruntime package, version 1.21.

[14] J. Pineau, P. Vincent-Lamarre, K. Sinha, V. Larivière, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and H. Larochelle. 2021. Improving reproducibility in machine learning research (a report from the NeurIPS 2019 reproducibility program). *Journal of Machine Learning Research* 22, 164 (2021), 1–20.

[15] D. R. Roberts, V. Bahn, S. Ciuti, M. S. Boyce, J. Elith, G. Guillera-Arroita, S. Hauenstein, J. J. Lahoz-Monfort, B. Schröder, W. Thuiller, D. I. Warton, B. A. Wintle, F. Hartig, and C. F. Dormann. 2017. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* 40, 8 (2017), 913–929. https://doi.org/10.1111/ecog.02881

[16] S. Timokhin, M. Sadrani, and C. Antoniou. 2020. Predicting venue popularity using crowd-sourced and passive sensor data. *Smart Cities* 3, 3 (2020), 818–841. https://doi.org/10.3390/smartcities3030042

[17] X. Zhou, D. Wang, and P. Krähenbühl. 2019. Objects as points. arXiv:1904.07850. https://doi.org/10.48550/arXiv.1904.07850

[18] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár. 2017. Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV '17)*. IEEE, 2999–3007. https://doi.org/10.1109/ICCV.2017.324

[19] V. Lempitsky and A. Zisserman. 2010. Learning to count objects in images. In *Advances in Neural Information Processing Systems 23 (NIPS 2010)*. 1324–1332.

[20] Y. Zhang, D. Zhou, S. Chen, S. Gao, and Y. Ma. 2016. Single-image crowd counting via multi-column convolutional neural network. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR '16)*. IEEE, 589–597. https://doi.org/10.1109/CVPR.2016.70

[21] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, and P. Abbeel. 2017. Domain randomization for transferring deep neural networks from simulation to the real world. In *Proceedings of the 2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 23–30. https://doi.org/10.1109/IROS.2017.8202133

[22] C. Pramerdorfer, J. Strohmayer, and M. Kampel. 2020. SDT: a synthetic multi-modal dataset for person detection and pose classification. In *Proceedings of the 2020 IEEE International Conference on Image Processing (ICIP)*. IEEE, 1611–1615. https://doi.org/10.1109/ICIP40778.2020.9191284

[23] A. Samal and H. R. Lone. 2025. Thermal vision: pioneering non-invasive temperature tracking in congested spaces. *Smart Health* 36 (2025), 100576. https://doi.org/10.1016/j.smhl.2025.100576. Dataset: https://osf.io/9z5a4/

[24] C. Palmero, A. Clapés, C. Bahnsen, A. Møgelmose, T. B. Moeslund, and S. Escalera. 2016. Multi-modal RGB–depth–thermal human body segmentation. *International Journal of Computer Vision* 118, 2 (2016), 217–239. https://doi.org/10.1007/s11263-016-0901-x. Dataset: https://vap.aau.dk/

[25] P. Aszkowski and M. Piechocki. 2022. Thermo Presence: the low-resolution thermal image dataset and occupancy detection using edge devices. In *Proceedings of the 3rd Polish Conference on Artificial Intelligence (PP-*

RAI 2022). Publishing House of Gdynia Maritime University, Gdynia, 49–52. Dataset: https://github.com/PUTvision/thermo-presence

[26] A. Kuzdeuov, M. Zakaryanov, A. Tleuliyev, and H. A. Varol. 2025. OpenThermalPose2: extending the open-source annotated thermal human pose dataset with more data, subjects, and poses. *IEEE Transactions on Biometrics, Behavior, and Identity Science* 7, 4 (2025), 902–913. https://doi.org/10.1109/TBIOM.2025.3575499. Dataset: https://huggingface.co/datasets/issai/OpenThermalPose2

[27] S. Zhu, T. Voigt, D. F. Perez-Ramirez, and J. Eriksson. 2021. A low-resolution infrared thermal dataset and potential privacy-preserving applications. In *Proceedings of the 19th ACM Conference on Embedded Networked Sensor Systems (SenSys '21)*. ACM, 552–555. https://doi.org/10.1145/3485730.3493692. Dataset: https://doi.org/10.5281/zenodo.5574233

[28] J. D. C. Little. 1961. A proof for the queuing formula: $L = \lambda W$. *Operations Research* 9, 3 (1961), 383–387. https://doi.org/10.1287/opre.9.3.383

[29] Teledyne FLIR. FLIR Lepton 3.5 radiometric thermal camera module. Product documentation, accessed October 2026.

[30] K. He, X. Zhang, S. Ren, and J. Sun. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR '16)*. IEEE, 770–778. https://doi.org/10.1109/CVPR.2016.90

[31] F. Yu and V. Koltun. 2016. Multi-scale context aggregation by dilated convolutions. In *International Conference on Learning Representations (ICLR 2016)*. arXiv:1511.07122.

[32] H. Law and J. Deng. 2018. CornerNet: detecting objects as paired keypoints. In *Computer Vision – ECCV 2018* (Lecture Notes in Computer Science). Springer, Cham, 765–781. https://doi.org/10.1007/978-3-030-01264-9_45

[33] Y. Li, X. Zhang, and D. Chen. 2018. CSRNet: dilated convolutional neural networks for understanding the highly congested scenes. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR '18)*. IEEE, 1091–1100. https://doi.org/10.1109/CVPR.2018.00120

[34] M. Piechocki, M. Kraft, T. Pajchrowski, P. Aszkowski, and D. Pieczyński. 2022. Efficient people counting in thermal images: the benchmark of resource-constrained hardware. *IEEE Access* 10 (2022), 124835–124847. https://doi.org/10.1109/ACCESS.2022.3225233

Appendix A: Choices made after seeing data

Hypotheses H1 to H3 were written after the preliminary findings of Section 3 and several of the analyses in Section 6 had been run; Section 7.2 is the first test fixed in advance. The persistence measurements of Section 3.2 cover 2026-09-01 to 09-08 and so overlap the three test days; they informed the decision to add a nowcast, while its weights were fitted on 09-01 to 09-05. The 30-day window of Section 6.1 was scored after the three-day result was known; its boundaries are the day after that test and the date of the export, and every model-served live reading between them is counted. The serving clamp, the one-point push that model mode gives rounded scores outside 25 to 65 toward the nearer end of the scale, the display map, the offset weight and every nowcast weight were chosen on a split earlier than the rows that report them. Weekly rows carried the weather of their collection moment, a leak cleared from 3,246,456 rows. Of rows recording “no event nearby”, 59.5% were schema defaults correlated with city (56.1% of the whole corpus), which is why Section 6.4 disables the events lookup for both arms. The persistence baselines, the two-way bootstrap, the age-mix standardisation, the readings

whose value changed and the test-window stickiness figures were computed on 13 October, after the first version of this paper; rule (b) was identified on the 30-day window and has no out-of-sample figure yet. The three held-out cities scored several model versions before v2.6.0 passed, as Section 6.2 states.